

Vel Tech High Tech

Dr.Rangarajan Dr.Sakunthala Engineering College
(AN AUTONOMOUS INSTITUTION)

**(Approved by AICTE, New Delhi & Affiliated to Anna University)
No.60, Avadi – Vel Tech Road, Chennai-600 062.**

21HC76P – FUNDAMENTALS OF CLOUD COMPUTING LABORATORY



NAME :

REGISTER NO :

ROLLNO :

BRANCH :

YEAR :

SEMESTER :

Vel Tech High Tech

Dr.Rangarajan Dr.Sakunthala Engineering College

(AN AUTONOMOUS INSTITUTION)

(Approved by AICTE, New Delhi & Affiliated to Anna University)

No.60, Avadi – Vel Tech Road, Chennai- 600 062.



BONAFIDE CERTIFICATE

Name: _____

Year: _____ Semester: _____ Branch: **B.E – Computer Science and Engineering**

University Register.....College Roll No:

Certified that this is the bonafide record of work done by the above student **21HC76P–**
FUNDAMENTALS OF CLOUD COMPUTING LABORATORY during the academic year 2025 -2026.

Signature of Course In charge

Signature of Head of the Department

Submitted for the University Practical Examination held on..... at **VEL TECH HIGH
TECH Dr.RANGARAJAN Dr.SAKUNTHALA ENGINEERING COLLEGE, No.60, AVADI
– VEL TECH ROAD, AVADI, CHENNAI-600062.**

Signature of Examiners

Internal Examiner:

External Examiner:

Date:

INDEX

S.No.	Date	Name of the Experiment	Page No	Marks	Signature
1		Install Virtual box/VMware/ Equivalent open source cloud Workstation with different Flavors of Linux or Windows OS on top of windows 8 and above.	1		
2		Install a C compiler in the virtual machine created using a virtual box and execute Simple Programs	13		
3		Install Google App Engine. Create a hello world app and other simple web applications Using python/java.	24		
4		Use the GAE launcher to launch the web applications.	30		
5		Simulate a cloud scenario using Cloud Sim and run a scheduling algorithm that is not Present in Cloud Sim.	41		
6		Find a procedure to transfer the files from one virtual machine to another virtual machine.	45		
7		Install Hadoop single node cluster and run simple applications like word count.	50		
8		Creating and Executing Your First Container Using Docker.	65		
9		Run a Container from Docker Hub	68		
10		Create a virtualization in VMWARE or any equivalent Open Source Tool. Allocate memory and storage space as per requirement and install Guest OS on that VMWARE	71		
11		To Host a web page in EC2 using AWS.	74		
12		To Create Billing Alarm using AWS Account.	78		
CONTENT BEYOND THE SYLLABUS					
13		Word mapper program to demonstrate the use of map and reduce tasks	80		
14		Cloud storage in Drop Box	82		

VEL TECH HIGH TECH DR. RANGARAJAN DR. SAKUNTHALA

ENGINEERINGCOLLEGE

INSTITUTION VISION

Pursuit of excellence in technical education to create civic responsibility with competency.

INSTITUTION MISSION

1. To impart the attributes of global engineers to face industrial challenges with social relevance.
2. To indoctrinate as front runners through moral practices.
3. To attain the skills through life long learning.

DEPARTMENT VISION

To blend academic learning process and innovative ideas producing self-confident graduates withskills and knowledge to compete in the changing world.

DEPARTMENT MISSION

M1: To provide strong basic knowledge in Computer Science and Engineering.

M2: To motivate the students to function as a team from their innovative ideas in collaboration withthe industries.

M3: To enable the students to lead and adapt to the growing environment.

PROGRAM EDUCATIONAL OBJECTIVES (PEOs)

PEO 1:

Embark upon successful professional practice in Computer Science and Engineering, displaying supportive and leadership roles.

PEO 2:

Engage in professional projects requiring teamwork and making valuable contributions to design, development, and production in the practice of Computer Science and Engineering or application areas.

PEO 3:

Equip to adapt and grow with changes in technology and globalization, and topursue higher studies and research activities.

PEO 4:

Be capable of productive employment in the field of Computer Science andEngineering with competing technical expertise, good interpersonal skill.

PEO 5:

Utilize their broad educational experience, ethics, and professionalism to make apositive impact on their local and professional communities.

PROGRAMME SPECIFIC OUTCOMES (PSOs)

By the time of graduation, the undergraduate Computer Science and Engineeringstudents can have the ability of

PSO's	PROGRAMME SPECIFIC OUTCOMES (PSOs)
PSO1	Designing Computer/Electronic based components which would serve social environment.
PSO2	Applying the current and gained knowledge and modern techniquesnot only in the Computers but in all related fields.

PROGRAM OUTCOMES:

1. **Engineering knowledge:** Apply the knowledge of mathematics, science, engineering fundamentals, and an engineering specialization to the solution of complex engineering problems.
2. **Problem analysis:** Identify, formulate, review research literature, and analyze complex engineering problems reaching substantiated conclusions using first principles of mathematics, natural sciences, and engineering sciences.
3. **Design/development of solutions:** Design solutions for complex engineering problems and design system components or processes that meet the specified needs with appropriate consideration for the public health and safety, and the cultural, societal, and environmental considerations.
4. **Conduct investigations of complex problems:** Use research-based knowledge and research methods including design of experiments, analysis and interpretation of data, and synthesis of the information to provide valid conclusions.
5. **Modern tool usage:** Create, select, and apply appropriate techniques, resources, and modern engineering and IT tools including prediction and modeling to complex engineering activities with an understanding of the limitations.
6. **The engineer and society:** Apply reasoning informed by the contextual knowledge to assess societal, health, safety, legal and cultural issues and the consequent responsibilities relevant to the professional engineering practice.
7. **Environment and sustainability:** Understand the impact of the professional engineering solutions in societal and environmental contexts, and demonstrate the knowledge of, and need for sustainable development.
8. **Ethics:** Apply ethical principles and commit to professional ethics and responsibilities and norms of the engineering practice.
9. **Individual and team work:** Function effectively as an individual, and as a member or leader in diverse teams, and in multi disciplinary settings.
10. **Communication:** Communicate effectively on complex engineering activities with the engineering community and with society at large, such as, being able to comprehend and write effective reports and design documentation, make effective presentations, and give and receive clear instructions.
11. **Project management and finance:** Demonstrate knowledge and understanding of the engineering and management principles and apply these to one's own work, as a member and leader in a team, to manage projects and in multi disciplinary environments.
12. **Life Long Learning:** Recognize the need for, and have the preparation and ability to engage in independent and life-long learning in the broadest context of technological change.

VEL TECH HIGTECH Dr. RANGARAJAN Dr. SAKUNTHALA
ENGINEERING COLLEGE

COURSE OBJECTIVES:

The student should be made to

- To develop Web applications in Cloud
- To learn the design and development process involved in creating a cloud based application
- To learn to implement and use parallel programming using Hadoop.
- To Run a Container from Docker Hub
- To Simulate Cloud Scenario using Cloud sim.
- To create virtualization in VMWARE or any equivalent Open Source Tool.

COURSE OUTCOMES:

CO. No.	Course Outcomes	Blooms level
On successful completion of this course, students will be able to		
C706. 1	Configure various virtualization tools such as Virtual Box, VMware workstation.	K2
C706. 2	Design and deploy a web application in a PaaS environment.	K3
C706. 3	Learn how to simulate a cloud environment to implement new schedulers.	K2
C706. 4	Install and use a generic cloud environment that can be used as a private cloud.	K3
C706. 5	Manipulate large data sets in a parallel environment	K3

COURSE OUTCOMES MAPPING WITH PROGRAM OUTCOMES AND PROGRAM SPECIFIC OUTCOMES

CO No.	PO-1	PO-2	PO-3	PO-4	PO-5	PO-6	PO-7	PO-8	PO-9	PO-10	PO-11	PO-12	PSO-1	PSO-2
C706. 1	2	1	1	1	1	-	-	-	-	-	-	2	2	2
C706. 2	2	1	1	1	1	-	-	-	-	-	-	2	2	2
C706. 3	2	1	1	1	1	-	-	-	-	-	-	2	2	2
C706. 4	3	2	2	1	2	-	-	-	-	-	-	2	3	2
C706. 5	2	1	1	1	1	-	-	-	-	-	-	2	2	2

Note:1:Slight,2: Moderate,3:Substantial

EXP.No:1	INSTALL VIRTUAL BOX/VMWARE EQUIVALENT OPEN-SOURCE CLOUD WORK STATION WITH DIFFERENT FLAVOURS OF LINUX OR WINDOWS OS ON TOP OF WINDOWS 8 AND ABOVE.
DATE:	

Aim:

PROCEDURE TO INSTALL

Step 1- Download Link

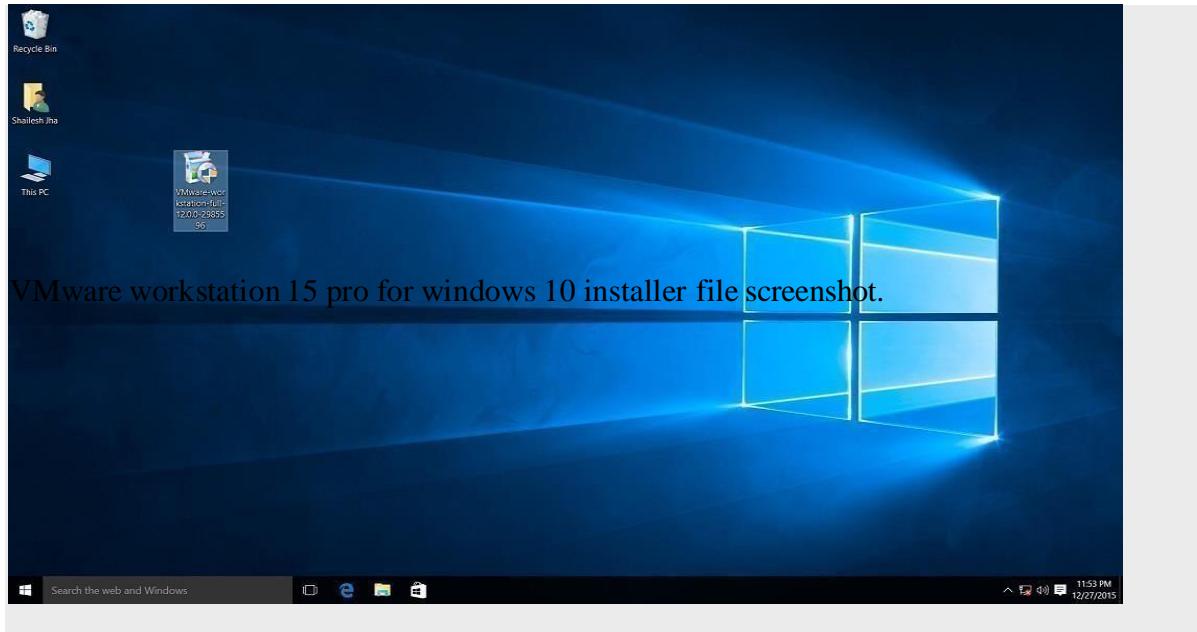
Link for downloading the software is <https://www.vmware.com/products/workstation-pro/workstation-pro-evaluation.html>. Download the software for windows. Good thing is that there is no sign up process. Click and download begins. Software is around 541MB.

Step 2- Download the installer file

It should probably be in the download folder by default, if you have not changed the settings in your browser. File name should be something like VMware-workstation-full-15.5.1-15018445.exe. This file name can change depending on the version of the software currently available for download. But for now, till the next version is available, they will all be VMware Workstation 15 Pro.

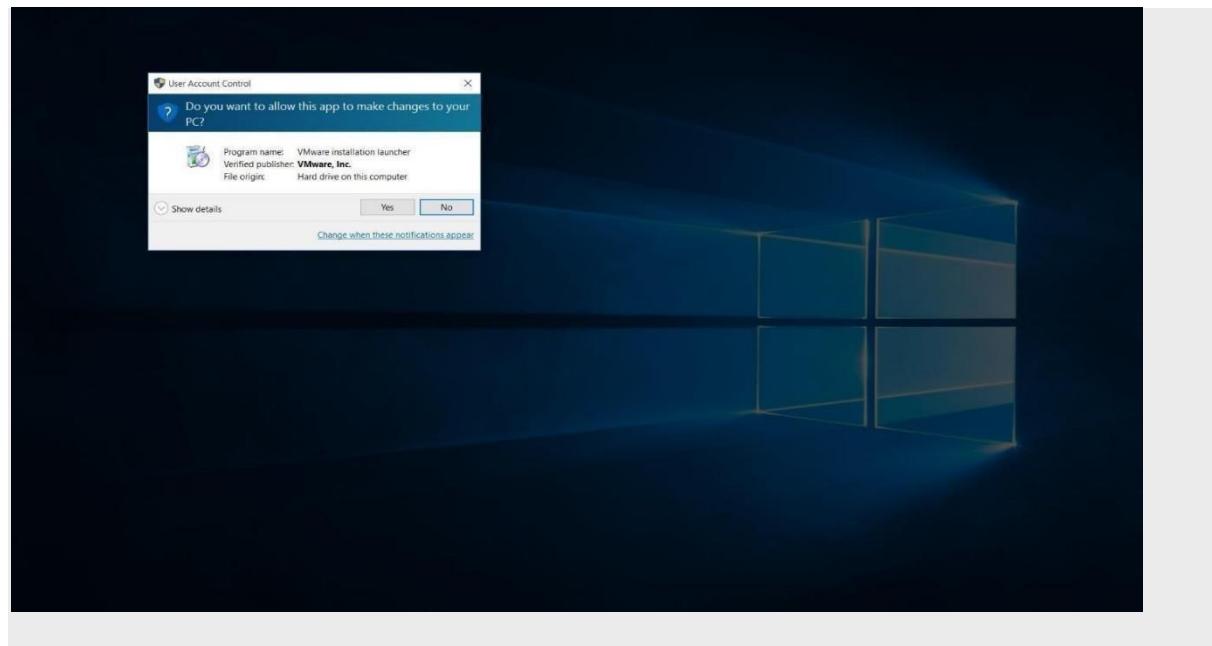
Step 3- Locate the downloaded installer file

For demonstration purpose, I have placed the downloaded installer on my desktop. Find the installer on your system and double click to launch the application.

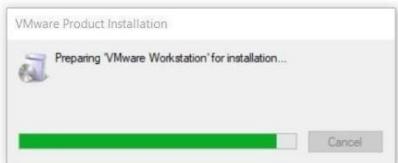
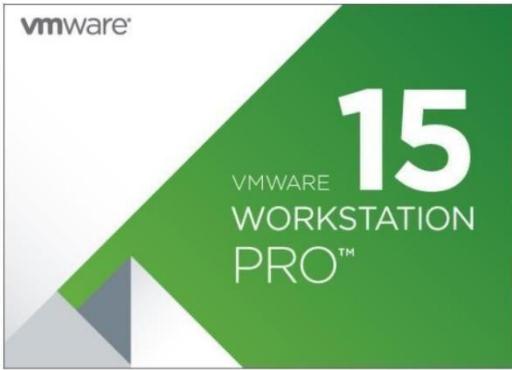


Step 4- User Access Control (UAC) Warning

Now you should see User Access Control (UAC) dialog box. Click yes to continue.



Initial Splash screen will appear. Wait for the process to complete.



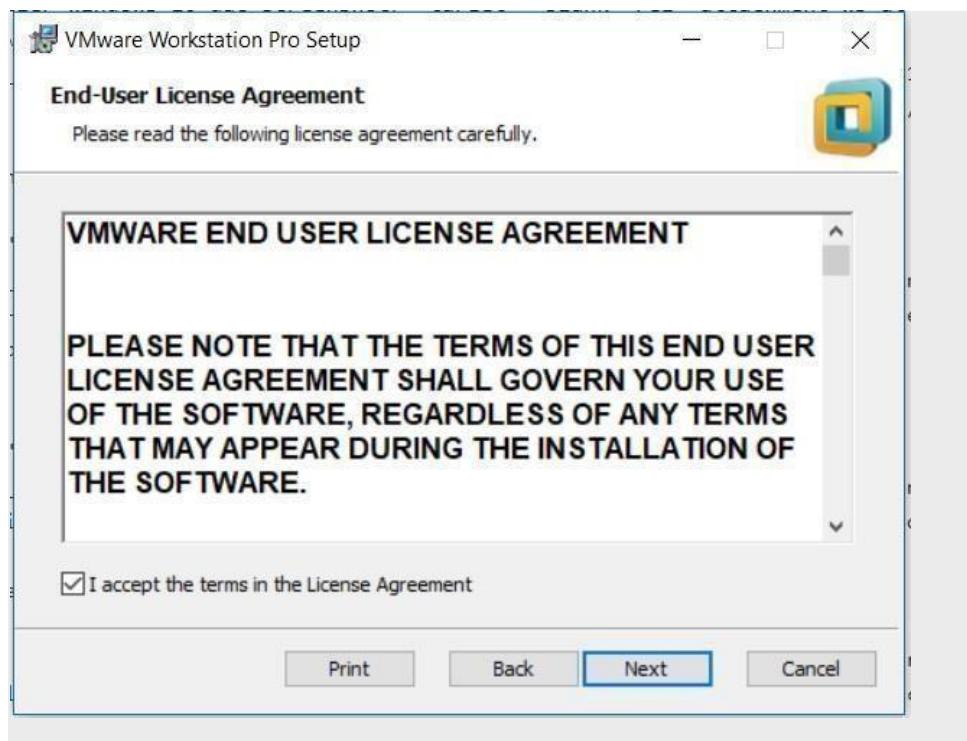
Step 5- VMware Workstation Setup wizard

Now you will see VMware Workstation setup wizard dialog box. Click next to continue.



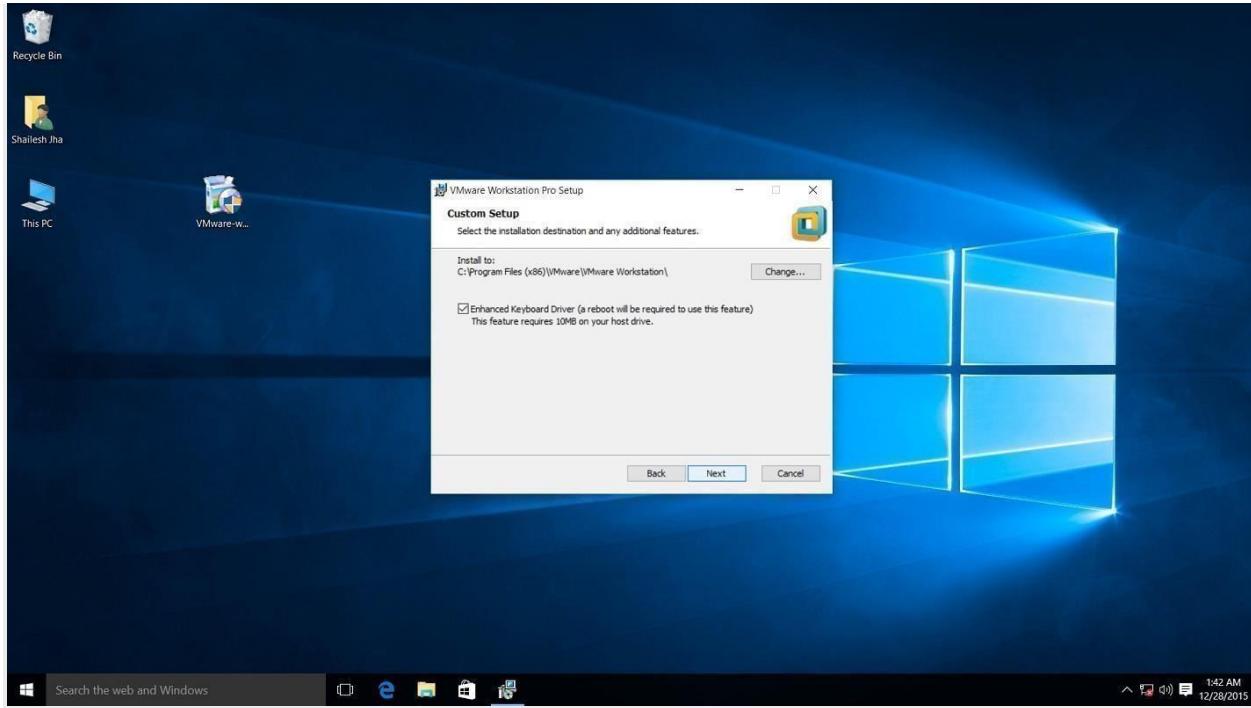
Step 6- End User Licence Agreement

This time you should see End User Licence Agreement dialog box. Check “I accept the terms in the Licence Agreement” box and press next to continue.



Step 7 - Custom Setup options

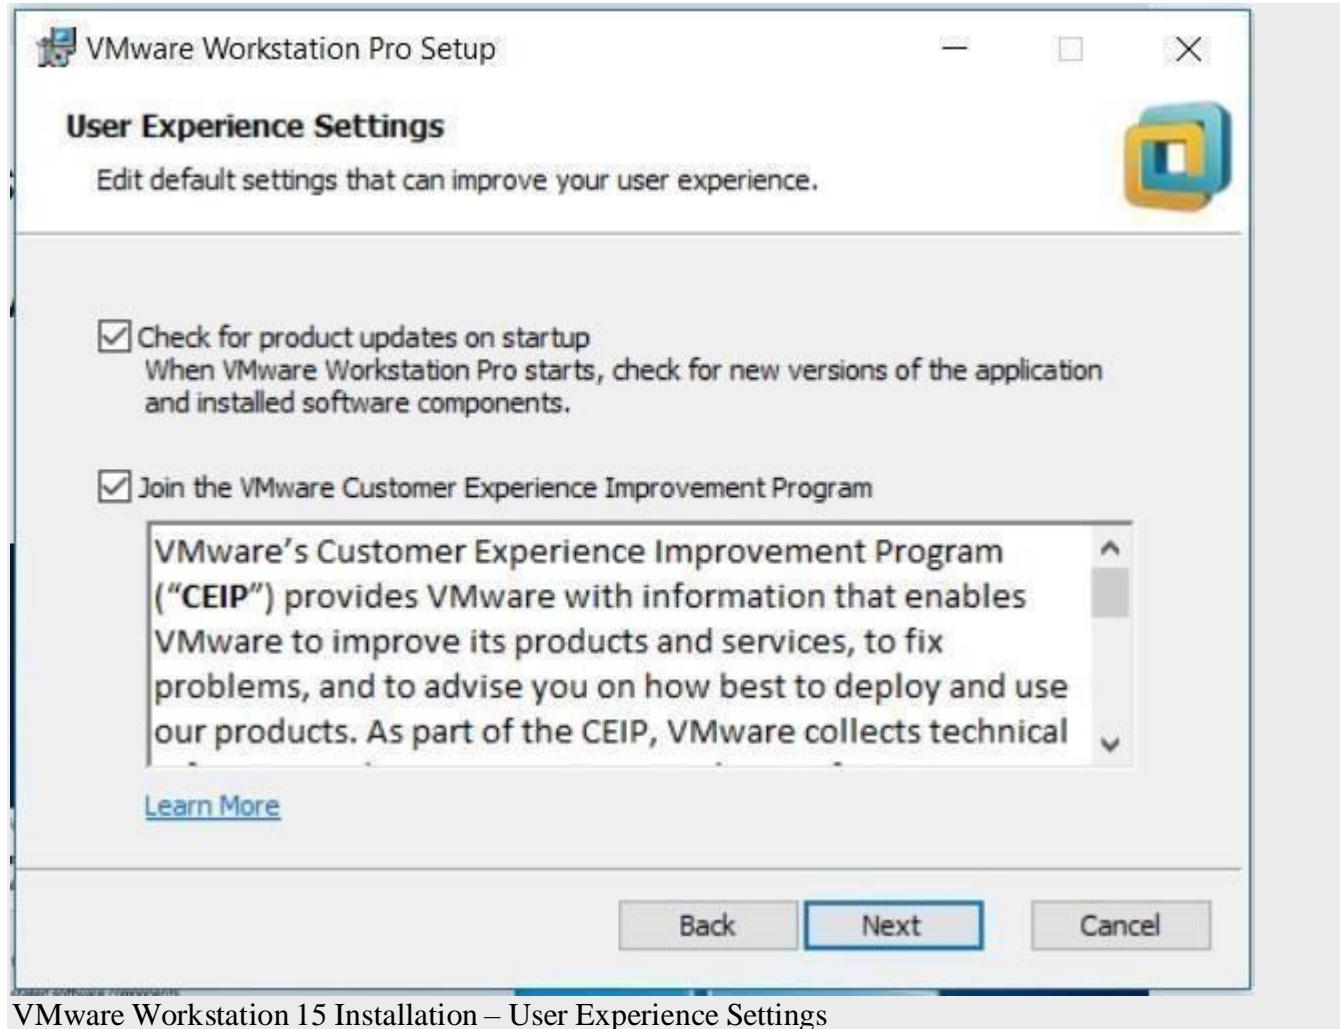
Select the folder in which you would like to install the application. There is no harm in leaving the defaults as it is. Also select Enhanced Keyboard Driver check box.



VMware Workstation 15 Pro installation – select installation folder

Step 8- User Experience Settings

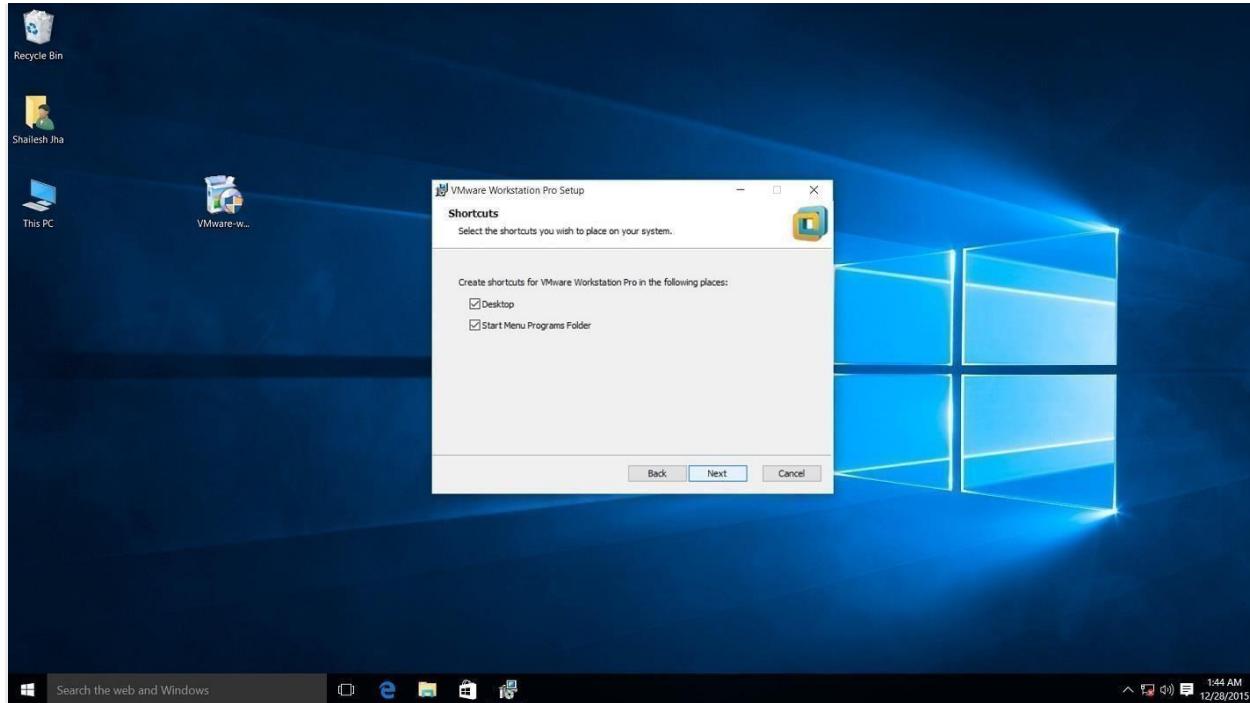
Next you are asked to select “Check for Updates” and “Help improve VMware Workstation Pro”. Do as you wish. I normally leave it to defaults that is unchecked.



VMware Workstation 15 Installation – User Experience Settings

Step 9- Application Shortcuts preference

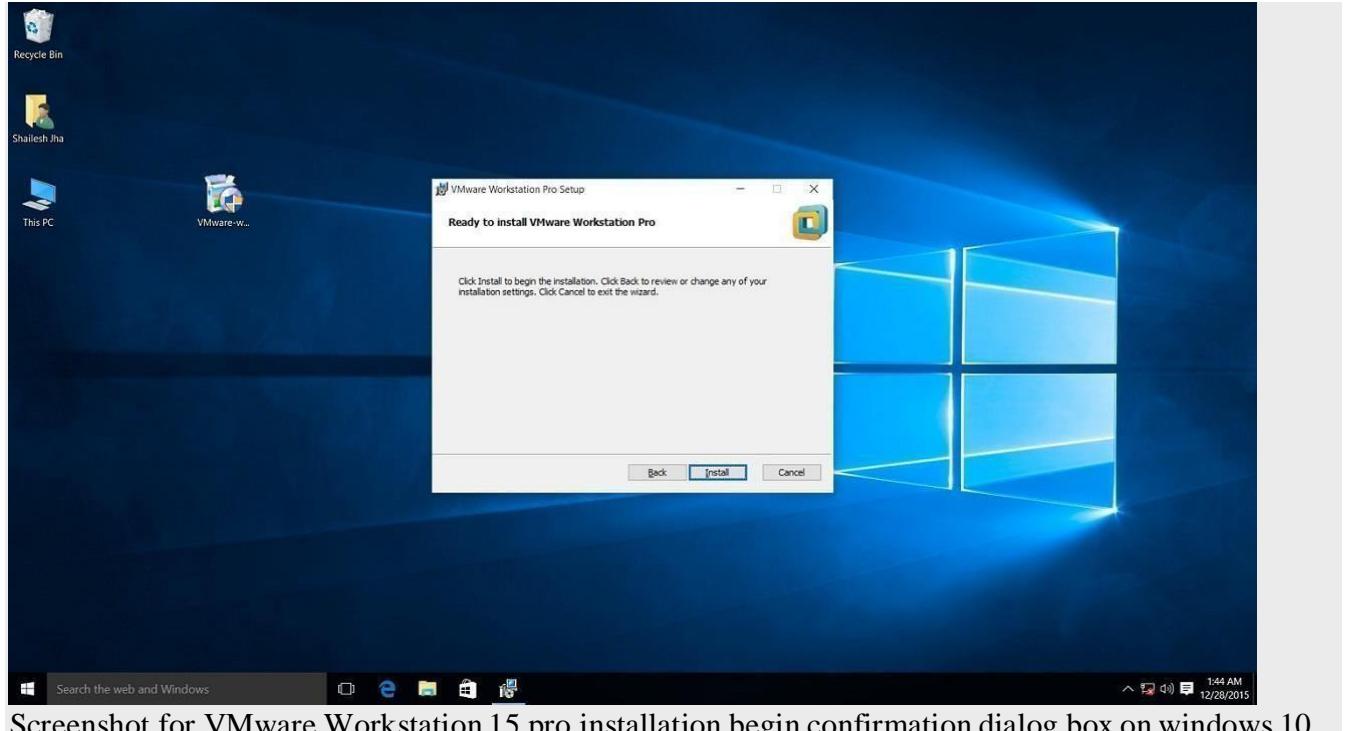
Next step is to select the place you want the shortcut icons to be placed on your system to launch the application. Please select both the options, desktop and start menu and click next.



VMware workstation 15 pro installation shortcut selection checkbox screenshot.

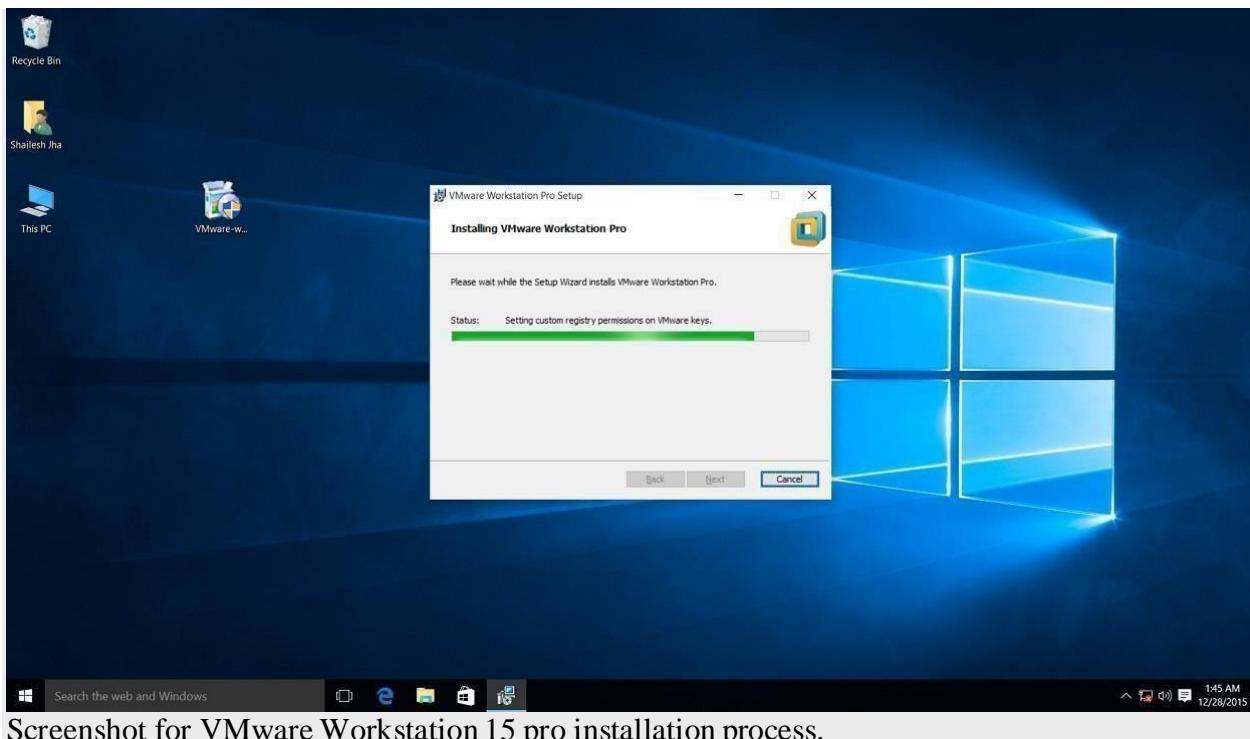
Step 10- Installation begins

Now you see the begin installation dialog box. Click install to start the installation process.



Screenshot for VMware Workstation 15 pro installation begin confirmation dialog box on windows 10.

Below screenshot shows Installation in progress. Wait for this to complete.



Screenshot for VMware Workstation 15 pro installation process.

At the end you will see installation complete dialog box. Click finish and you are done with the installation process. You may be asked to restart your computer. Click on Yes to restart.



Step 11- Launch VMware Workstation

After the installation completes, you should see VMware Workstation icon on the desktop. Double click on it to launch the application.



Screenshot for VMware Workstation 15 Pro icon on windows 10 desktop.

Step 12- Licence Key

If you see the dialog box asking for licence key, click on trial or enter the licence key. Then what you have is the VMware Workstation 15 Pro running on your windows 10 desktop. If don't have the licence key, you will have 30 day trial.



VMware Workstation 15 Pro home screen

Step 13- At some point if you decide to buy

At some point of time if you decide to buy the Licence key, you can enter the Licence key by going to **Help->Enter a Licence Key**. You can enter the 25 character licence key in the dialog box shown below and click OK. Now you have the licence version of the software.

Result:

EXP NO: 2
DATE:

**INSTALL A C COMPILER IN THE VIRTUAL MACHINE CREATED USING A VIRTUAL
BOX AND EXECUTE SIMPLE PROGRAMS**

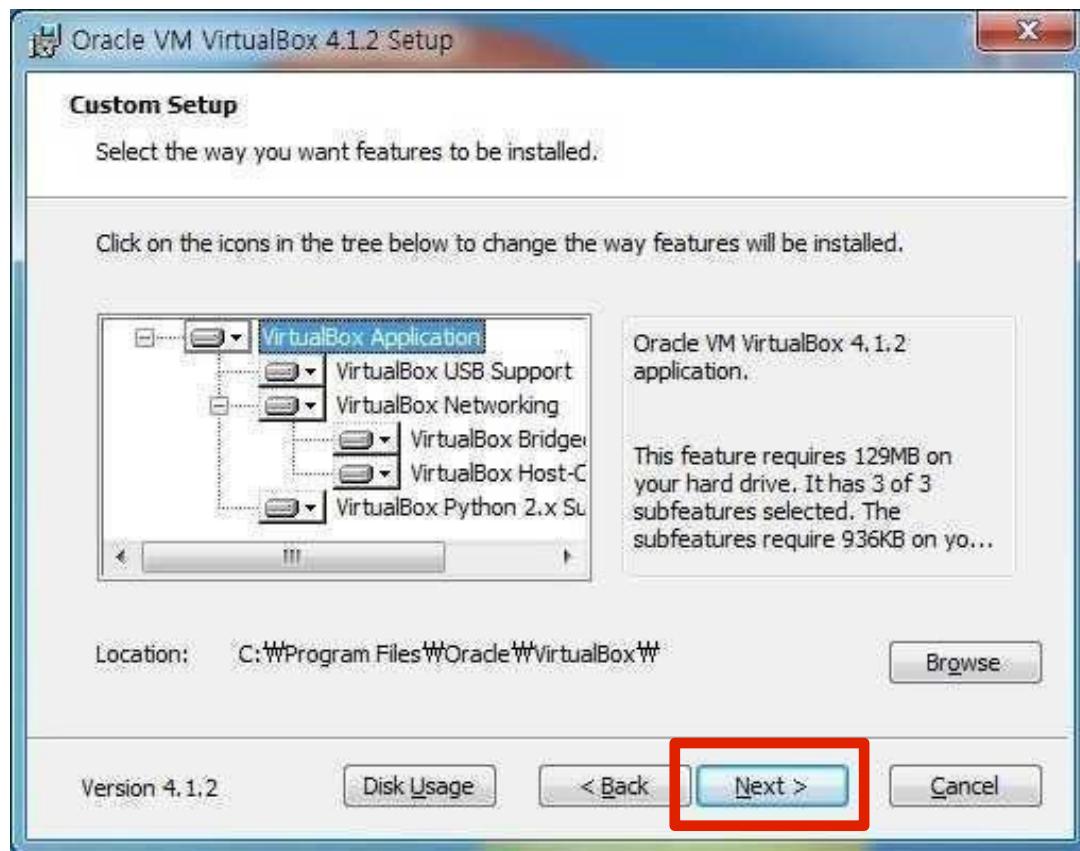
Aim:

1. A. Install VirtualBox

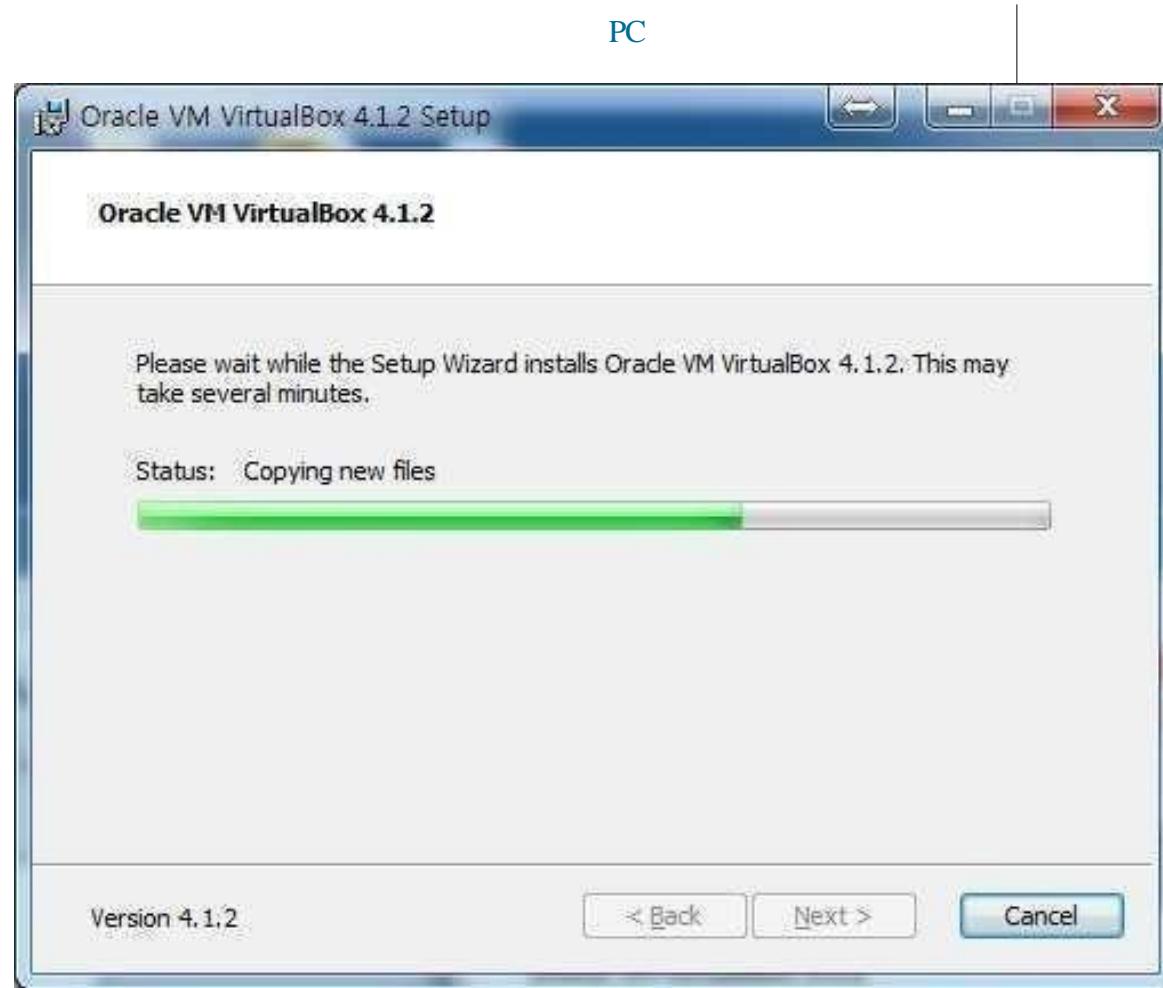
I. Visit <http://www.virtualbox.org/wiki/downloads>

- 2.** Download VirtualBox platform packages for yourOS

- 3.** Open the Installation Package by double clicking



4. Click continue and finish installing VirtualBox



5. When finished installation, close the window.

2. Download Linux

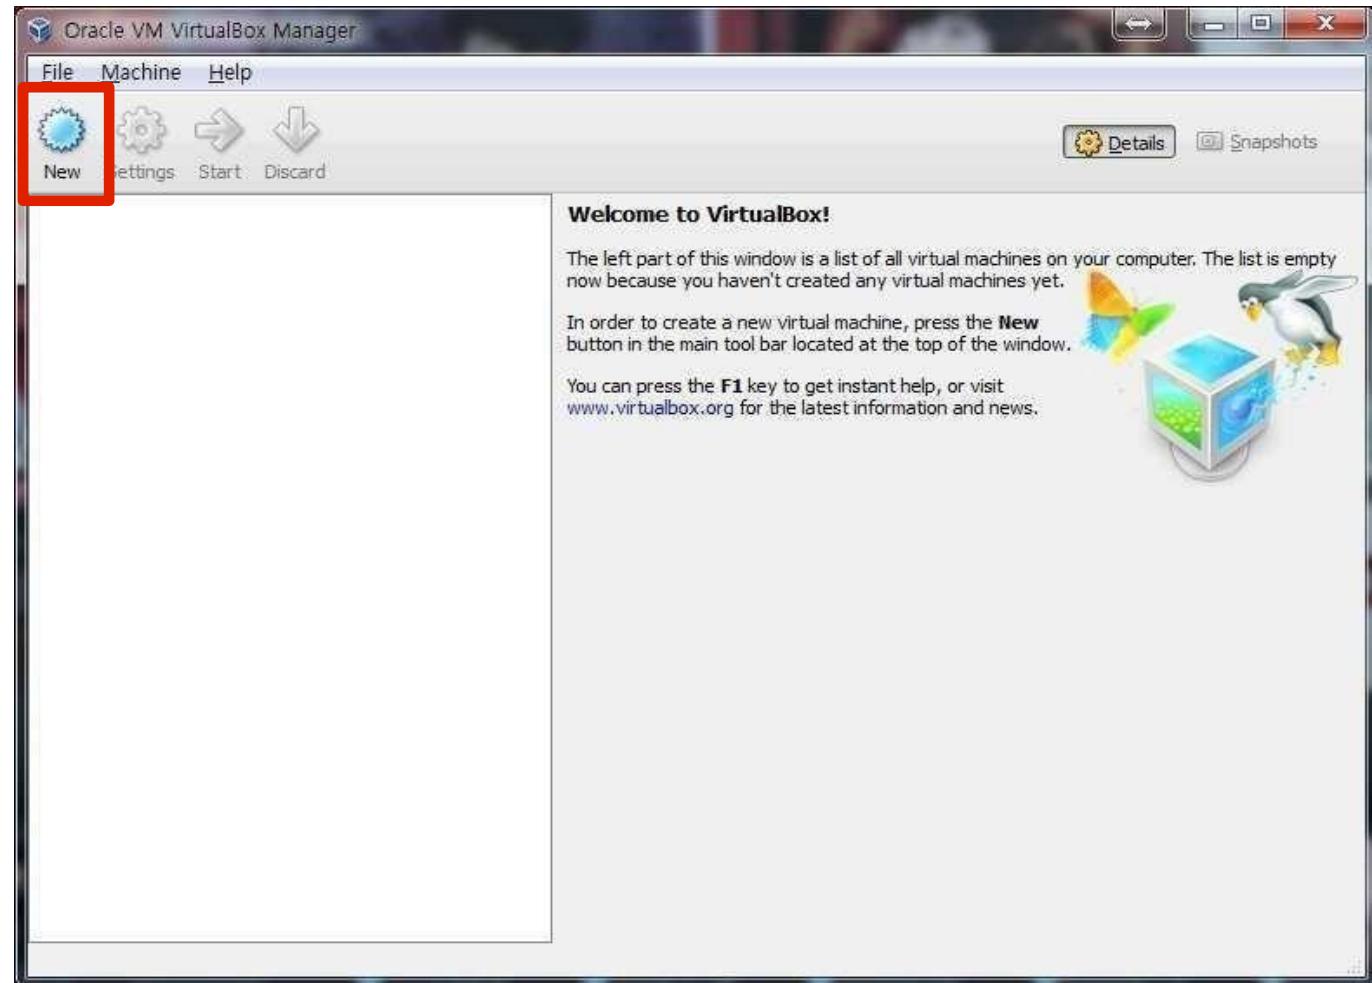
- Visit the page <http://www.ubuntu.com/download/ubuntu/download>
- Choose the Latest version of Ubuntu and 32-bit and click "Start Download"

The screenshot shows the Ubuntu download page. At the top, there are tabs for 'Download', 'Windows Installer', 'Alternative downloads', 'CDs', 'Upgrade', and '下载 Ubuntu'. A large orange button labeled '1 Download Ubuntu' is on the left. To its right, there's a section with text: 'Click the big orange button to download the latest version of Ubuntu. You will need to create a CD or USB stick to install Ubuntu.' Below this is another section: 'Our long-term support (LTS) releases are supported for three years on the desktop. Perfect for organisations that need more stability for larger deployments.' In the center, there's a 'Download options' box with two dropdown menus: 'Ubuntu 11.04 - Latest version' and '32-bit (recommended)', both highlighted with a red box. To the right of this is a large orange button with the text 'Download started' and 'Ubuntu 11.04 32-bit'. An arrow points from the word 'CLI' at the top right towards this button. At the bottom, there are sections for 'Additional options', 'If you're running Windows', and 'Other ways to get Ubuntu'.

I. Run VirtualBox by double-clicking the icon

2. Click “New” button on the top left corner

PC



3. Click “Continue” on the pop-up window

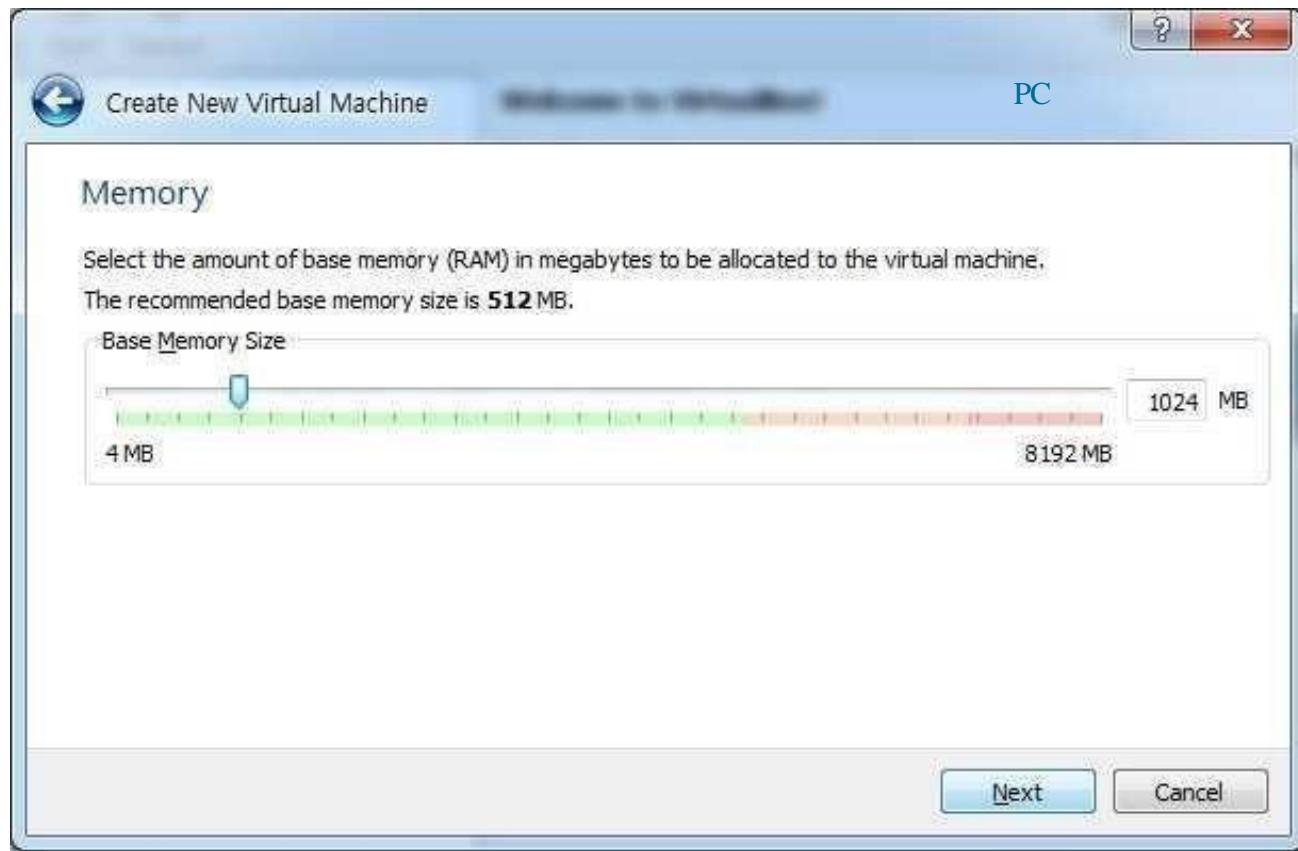
4. Type VM name, select “Linux” for the OS and choose “Ubuntu” for the version.



Page 19 of 71

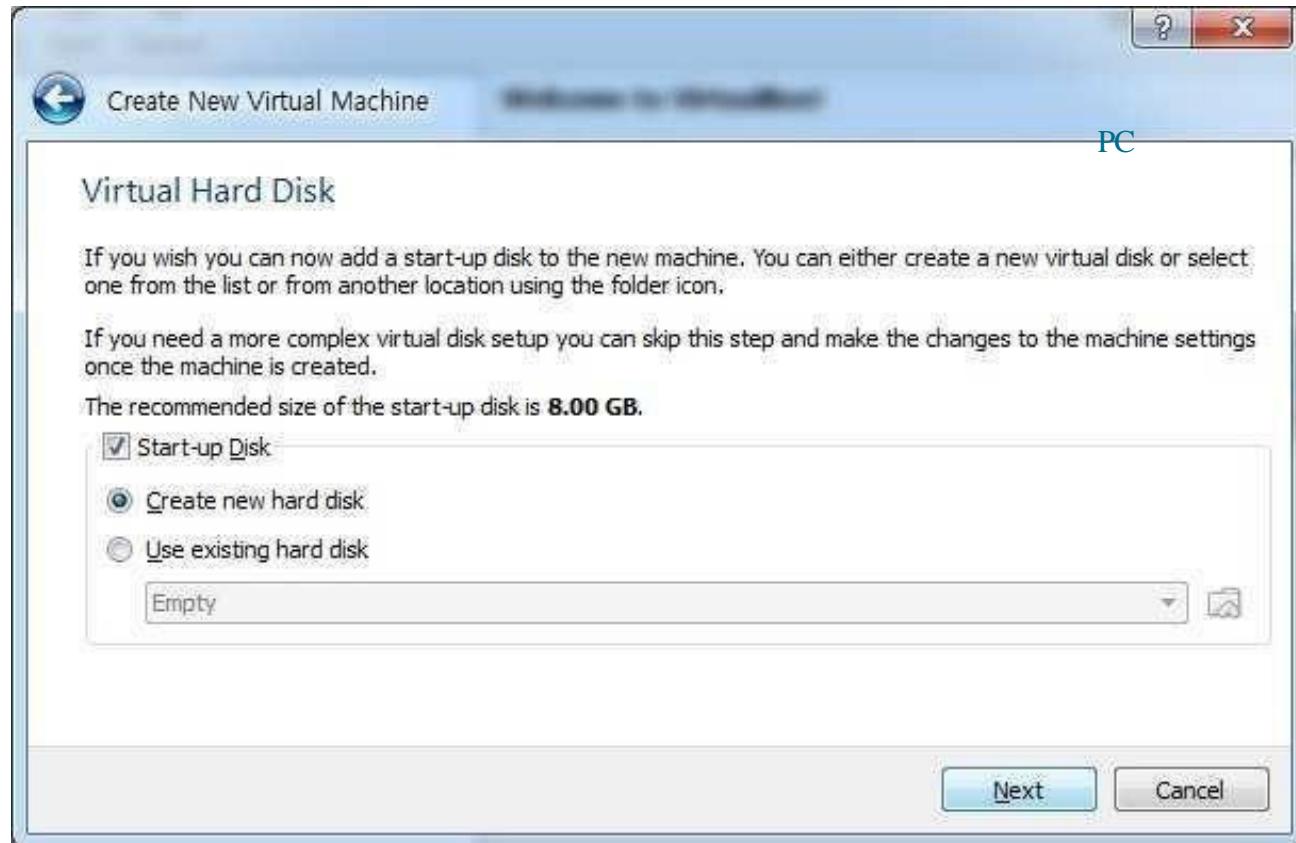
5. Choose the amount of memory to allocate (I suggest choosing between 512 MB to 1024MB)

6. Click Continue orNext



7. Choose create a new virtual harddisk

8. Click Continue orNext



9. Choose VDI (VirtualBox DiskImage) I0. Click Continue orNext



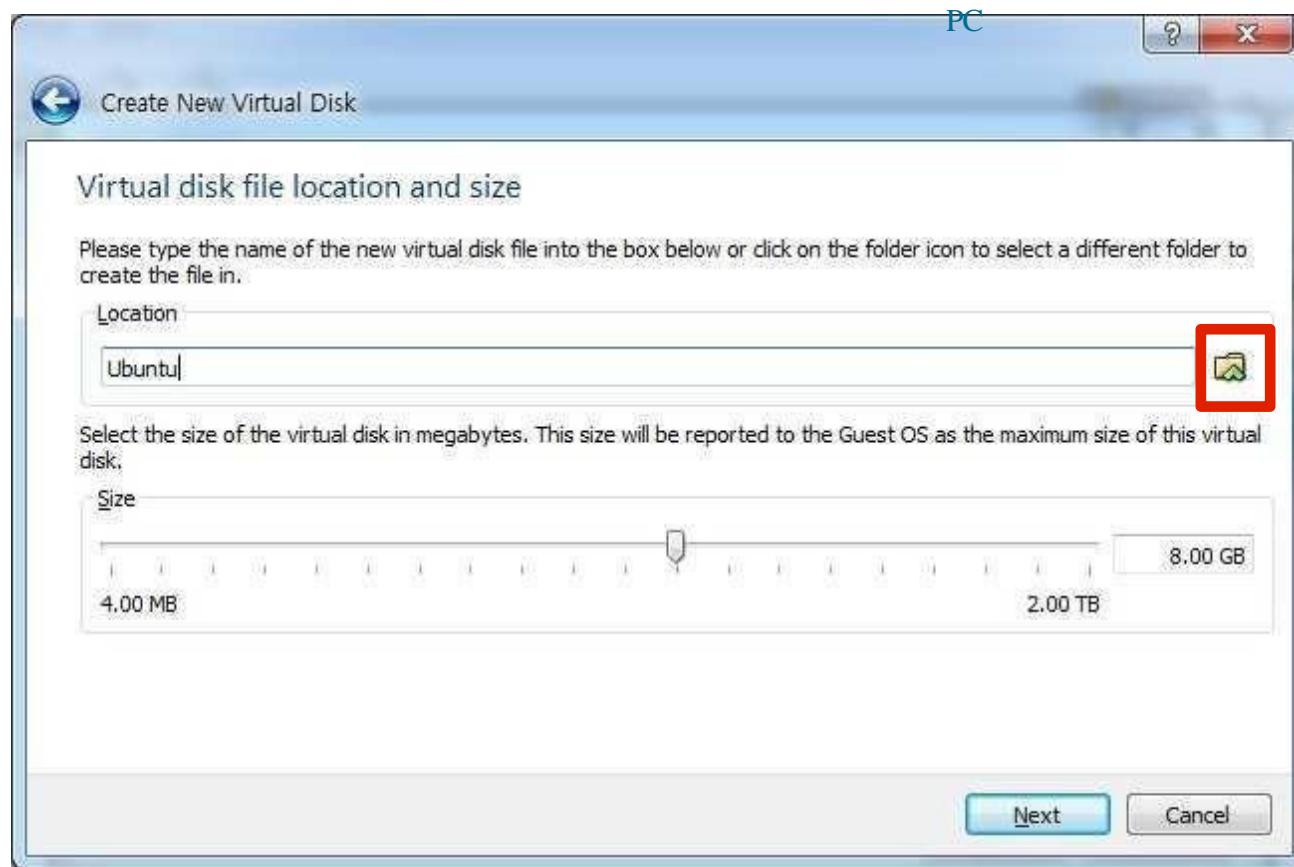
- II. Choose “DynamicallyAllocated” click continue.

This way, the size of your Virtual Hard Disk will grow as you use.

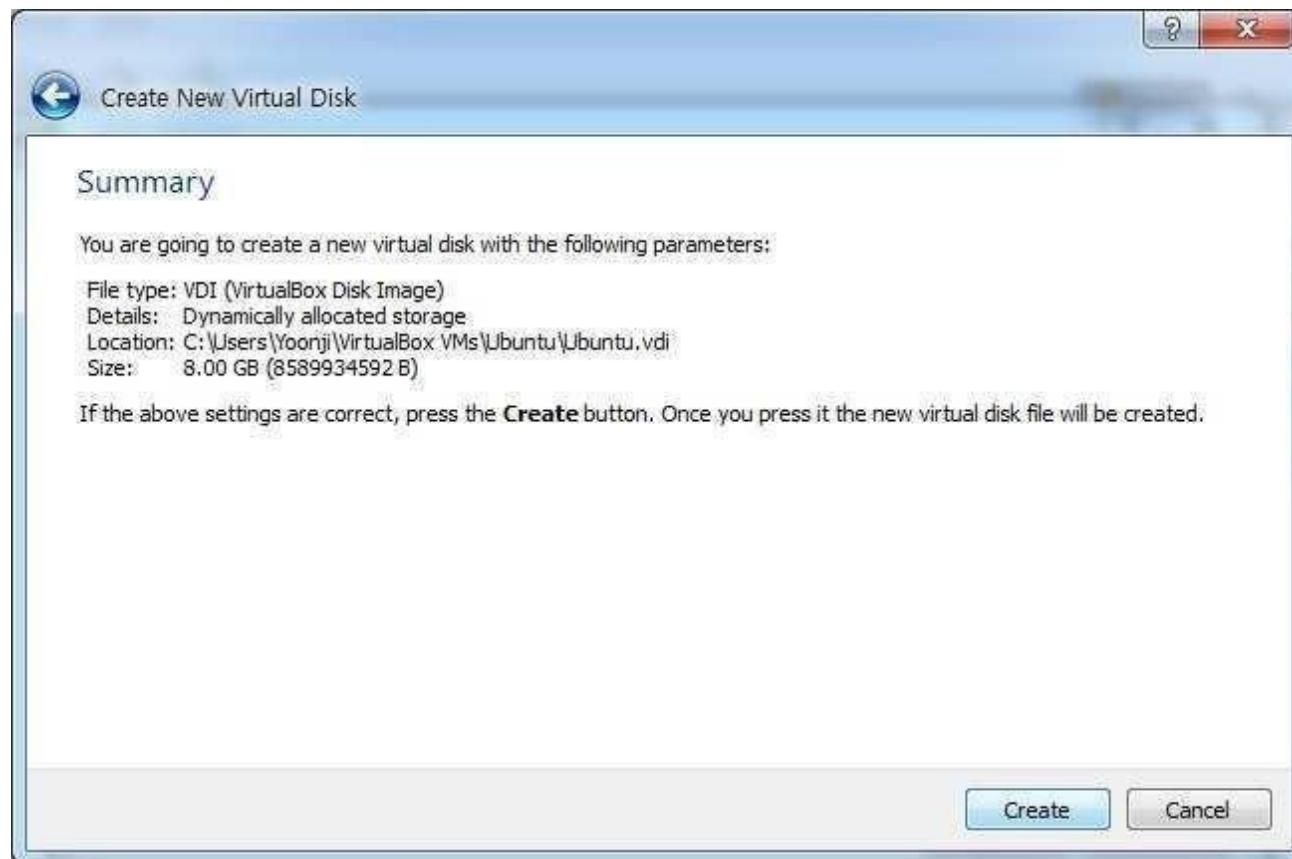


I2. Click the folder icon and choose the ubuntu iso file you downloaded.

I3. Select the size of the Virtual Disk (I recommend choosing 8 GB) and click continue

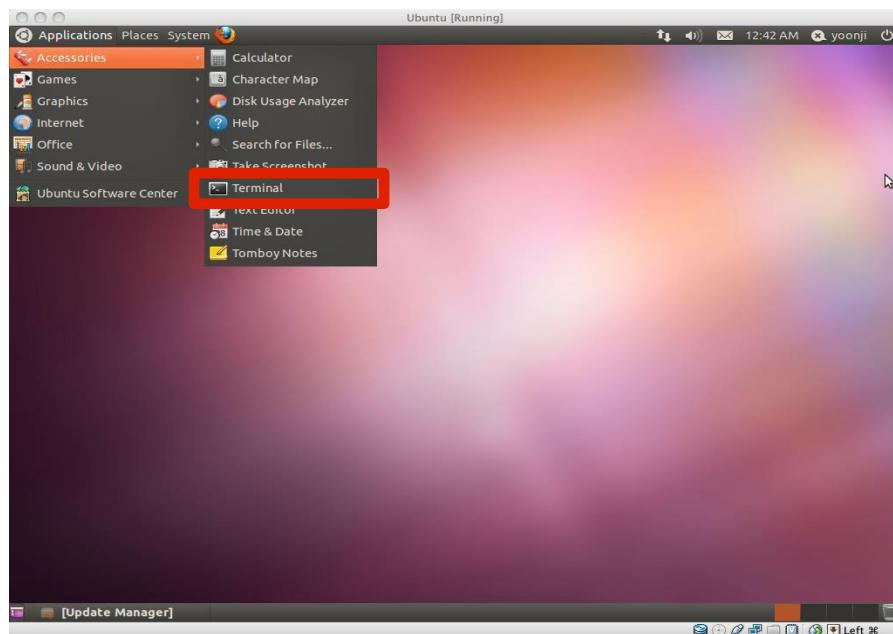


I4. Click Create



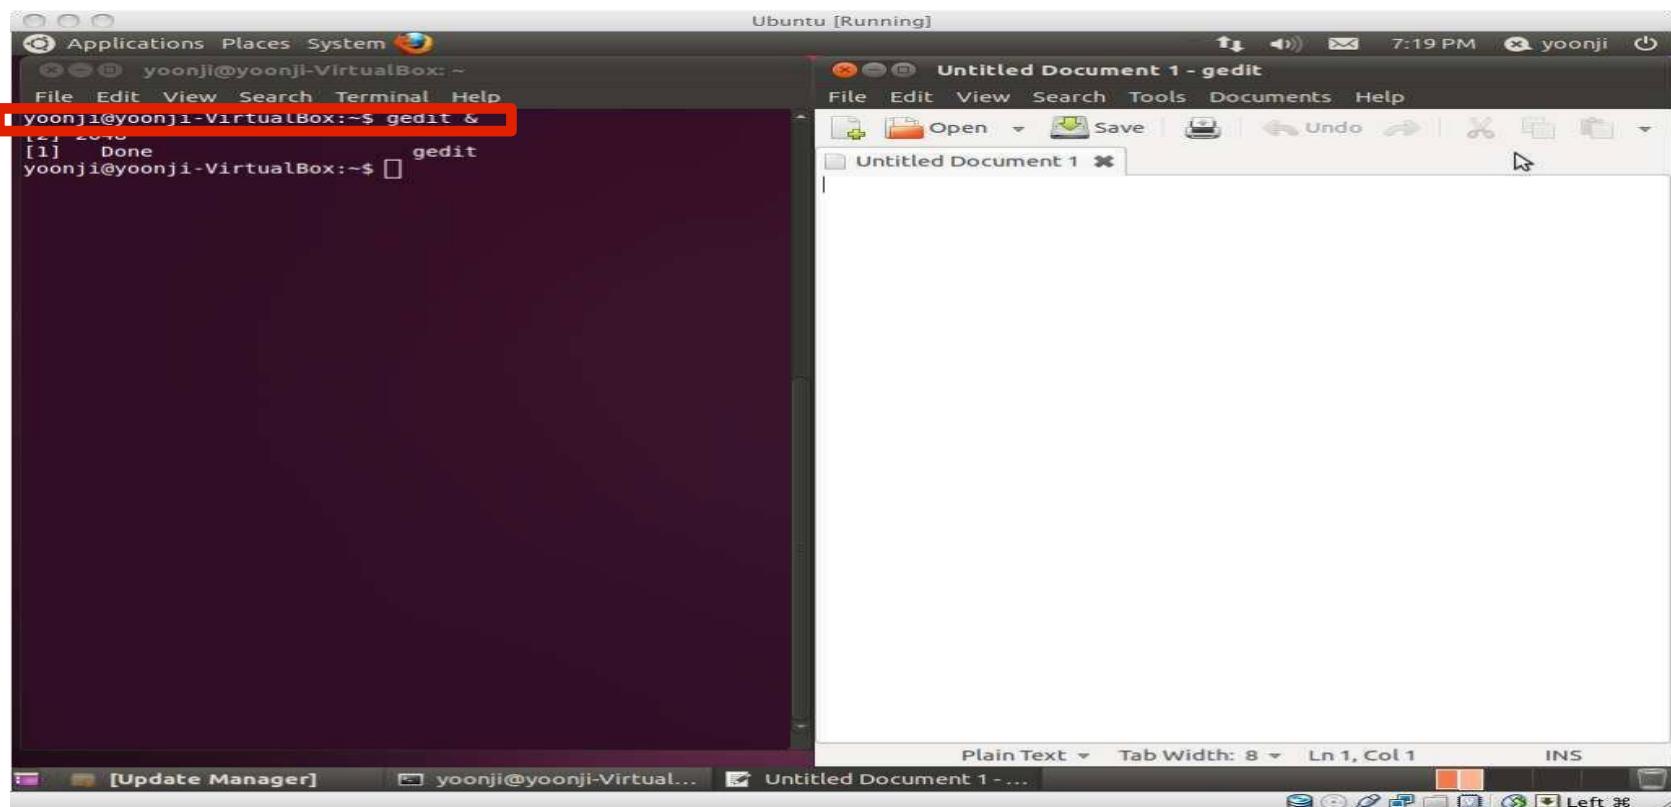
Exno.2.a : C Programming on Linux

I.O pen Terminal (Applications-Accessories-Terminal)



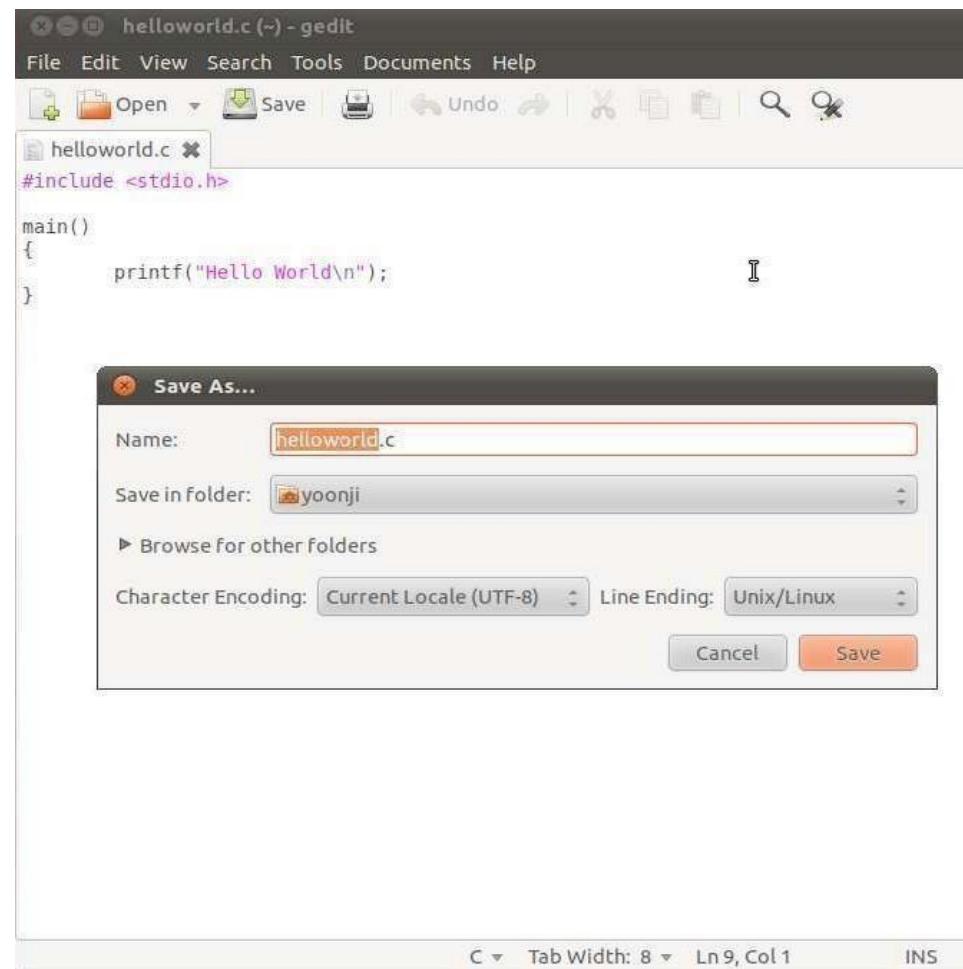
2. Open gedit by typing “gedit &” on terminal

(You can also use any other Text Editor application)



3. type program as

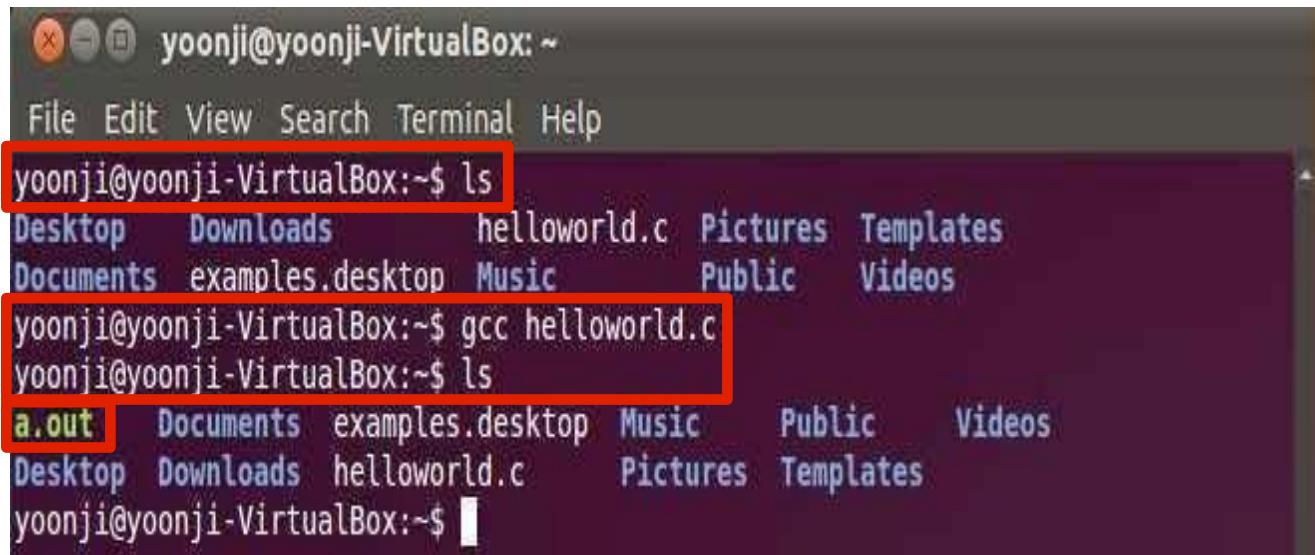
```
#include<stdio.h>
Void main()
{ printf("helloworld");
}
```



Save it as helloworld.c

Type "ls" on Terminal to see all files under current folder

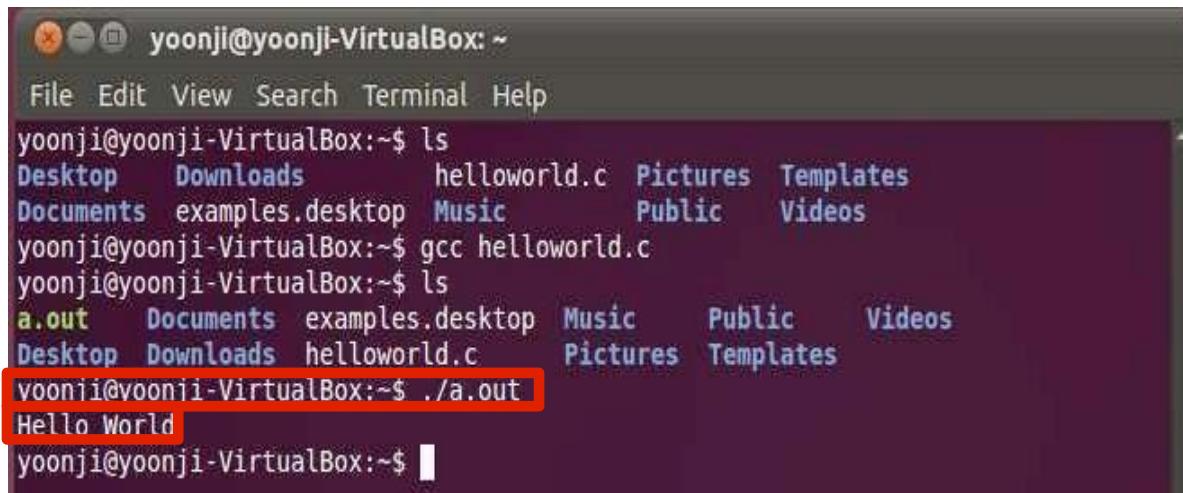
4. Confirm that "helloworld.c" is in the current directory. If not, type `cd DIRECTORY_PATH` to go to the directory that has "helloworld.c"
5. Type `gcc helloworld.c` to compile, and type `ls` to confirm that a new executable file "a.out" is created



6. Type `./a.out` on Terminal to run the program

7. If you see "HelloWorld" on the next line,
you just successfully ran your first C program!

8. Try other codes from “A Shotgun Introduction to C” on professor Edwards’s webpage. You can also find many C programming guides online. (just google it!)



The screenshot shows a terminal window titled "yoonji@yoonji-VirtualBox: ~". The window contains the following text:

```
File Edit View Search Terminal Help
yoonji@yoonji-VirtualBox:~$ ls
Desktop Downloads helloworld.c Pictures Templates
Documents examples.desktop Music Public Videos
yoonji@yoonji-VirtualBox:~$ gcc helloworld.c
yoonji@yoonji-VirtualBox:~$ ls
a.out Documents examples.desktop Music Public Videos
Desktop Downloads helloworld.c Pictures Templates
yoonji@yoonji-VirtualBox:~$ ./a.out
Hello World
yoonji@yoonji-VirtualBox:~$
```

The command `./a.out` and its output "Hello World" are highlighted with a red rectangle.

Result:

EXP NO: 3	INSTALL GOOGLE APP ENGINE : CREATE A HELLO WORLD APP AND OTHER SIMPLE WEB APPLICATIONS USING PYHTON / JAVA.
DATE:	

Aim:

Procedure:

- The App Engine SDK allows you to run Google App Engine Applications on your local computer. It simulates the run-time environment of the Google App Engine infrastructure.

Step1: To install python

Pre-Requisites: Python 2.5.4

If you don't already have Python 2.5.4 installed in your computer, download and Install Python 2.5.4 from:

<http://www.python.org/download/releases/2.5.4/>

Step 2: to install Google App Engine

Download and Install

You can download the Google App Engine SDK by going to:

<http://code.google.com/appengine/downloads.html> and download the appropriate installpackage.

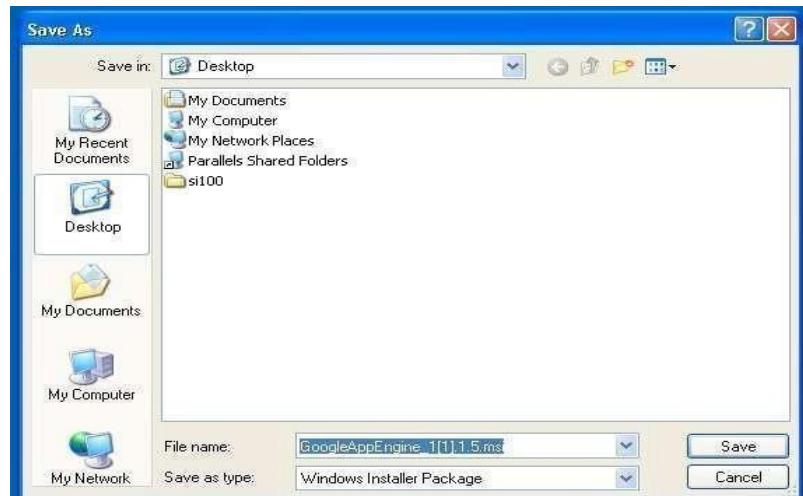
Download the Google App Engine SDK

Before downloading, please read the [Terms](#) that govern your use of the App Engine SDK.

Please note: The App Engine SDK is under **active development**, please keep this in mind as you explore its capabilities. See the [SDK Release Notes](#) for the information on the most recent changes to the App Engine SDK. If you discover any issues, please feel free to notify us via our [Issue Tracker](#).

Platform	Version	Package	Size	SHA1 Checksum
Windows	1.1.5 - 10/03/08	GoogleAppEngine_1.1.5.msi	2.5 MB	e974312b4aefc0b3873ff0d93eb4c525d5e88c30
Mac OS X	1.1.5 - 10/03/08	GoogleAppEngineLauncher-1.1.5.dmg	3.6 MB	f62208ac01c1b3e39796e58100d5f1b2f052d3e7
Linux/Other Platforms	1.1.5 - 10/03/08	google_appengine_1.1.5.zip	2.6 MB	cbb9ce817bdabf1c4f181d9544864e55ee253de1

Download the Windows installer – the simplest thing is to download it to your Desktop or another folder that you remember.



Double Click on the **GoogleApplicationEngine** installer.



Click through the installation wizard, and it should install the App Engine. If you do not have Python 2.5, it will install Python 2.5 as well.

Once the install is complete you can discard the downloaded installer

After installation Google app engine looks that,



Step 3: Making of the First Application

Now we need to create a simple application. We could use the “+” option to have the launcher make us an application – but instead we will do it by hand to get a better sense of what is going on.

Make a folder for your Google App Engine applications. I am going to make the Folder on my Desktop called “**apps**” – the path to this folder is:

C:\Documents and Settings\csev\Desktop\apps

And then make a sub-folder in **apps** called “**ae-01-trivial**”—the path to this folder would be:

C:\ Documents and Settings \csev\Desktop\apps\ae-01-trivial

Using a text editor such as JEdit (www.jedit.org), create a file called **app.yaml** in the **ae-01-trivial** folder with the following contents:

```
application: ae-01-
trivial
version: 1
runtime: python
api_version: 1
handlers:
- url: /.*
  script: index.py
```

Then create a file in the **ae-01-trivial** folder called **index.py** with three lines in it:

```
print 'Content-
Type:text/plain' print ''
print 'Hello there Chuck'
```

Step 4: Run the program

```
Administrator: Google Cloud Shell - dev_appserver.py app.yaml
Welcome to the Google Cloud SDK! Run "gcloud -h" to get the list of available commands.
---

C:\Program Files (x86)\Google\Cloud SDK>dev_appserver.py app.yaml

Updates are available for some Cloud SDK components. To install them,
please run:
  $ gcloud components update

This action requires the installation of components: [app-engine-
python, cloud-datastore-emulator]

Your current Cloud SDK version is: 316.0.0
Installing components from version: 316.0.0

  These components will be installed.
  +-----+-----+-----+
  | Name | Version | Size |
  +-----+-----+-----+
  | Cloud Datastore Emulator | 2.1.0 | 18.4 MiB |
  | gRPC python library | 1.20.0 | 1.5 MiB |
  | gRPC python library | 1.20.0 | 1.5 MiB |
  | gcloud app Python Extensions | 1.9.91 | 6.1 MiB |
  +-----+-----+-----+

For the latest full release notes, please visit:
  https://cloud.google.com/sdk/release_notes

Do you want to continue (Y/n)? y

  Creating update staging area
```

Output:

Once you have selected your application and press **Run**. After a few moments your application will start and the launcher will show a little green icon next to your application. Then press **Browse** to open a browser pointing at your application which is running at **http://localhost:8080/**

Paste **http://localhost:8080** into your browser and you should see your application as follows:



Result:

EXP NO: 4	USE THE GAE LAUNCHER TO LAUNGH THE WEB
DATE:	APPLICATIONS

Aim:

Procedure:

Step 1. Download the basic housekeeping stuff

No matter what platform you build products on, there is always some housekeeping stuff you need to put in place before you can hit the ground running. And deploying apps within the Google App Engine is no exception.

1. Download [Python 2.7](#)

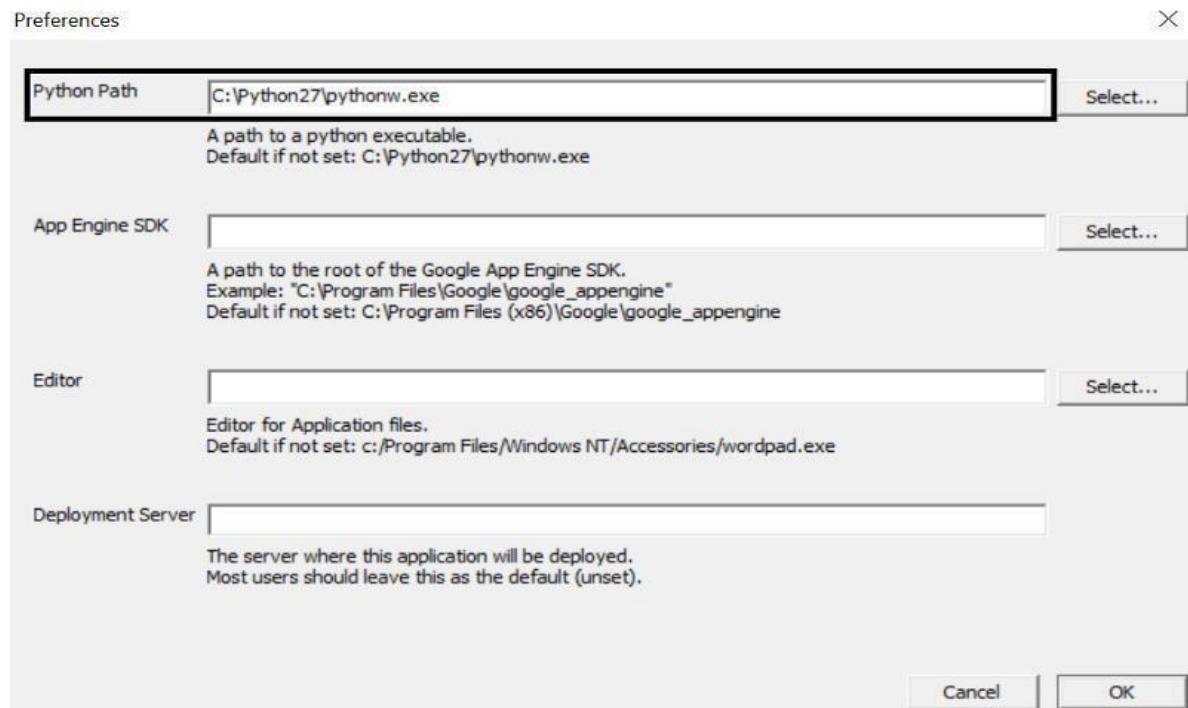
As of when this article was written, the Google App Engine [standard environment supports Python only upto version 2.7](#). However, it is only a matter of time before support for Python 3.x is added. You can check the App Engine docs for the latest info.

2. Download [Google Cloud SDK](#)

This will allow you to fork apps onto your local machine, make changes (edit and develop the app), and deploy your app back to the cloud.

3. Set the Python path in the Google App Engine launcher

After downloading the SDK, launch the App Engine launcher, go to Edit -> Preferences and make sure you set the path for where you installed Python in step 1 above.



Set the Python path in Google App Engine launcher

That's all you need. Your local machine should now be ready to build webapps.

Step 2. App Engine sign-up

This is often the most confusing part of the entire setup. Things you should know when you sign-up:

1. Currently, App Engine offers a free trial for one year.
2. The trial includes \$300 of credit that can be used during the one year trial period.
3. You will need to add a credit card to sign-up (for verification purposes).
4. You will not be charged during the sign-up process.
5. You will not be charged during the trial period as long as you do not cross the credit limit offered.

Here are the steps you need to follow to sign-up:

1. Go to the [Google Cloud](#) landing page
2. Follow the sign-up process and go to your App Engine dashboard

Most of the hard work is complete after a successful sign-up.

Step 3. Create a new project

The next step is to create a new Python project that you can work on. Follow the screenshots below to create a new project.

Launch the new project wizard.

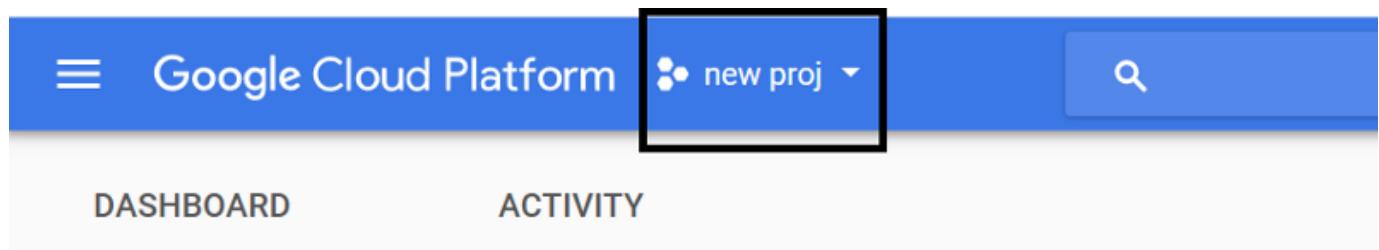


Image courtesy.

<https://console.cloud.google.com/home>

Select

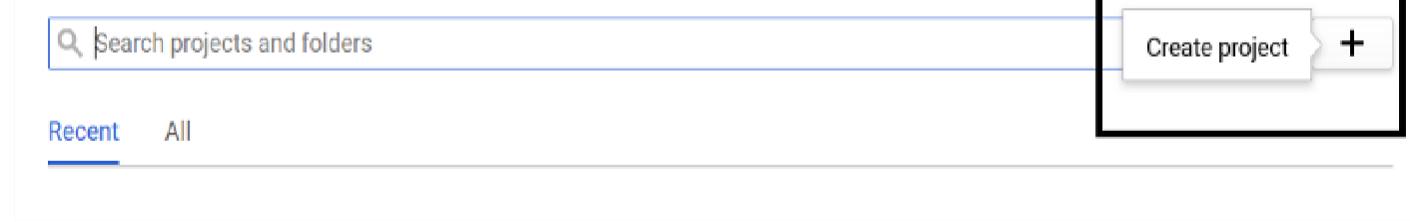
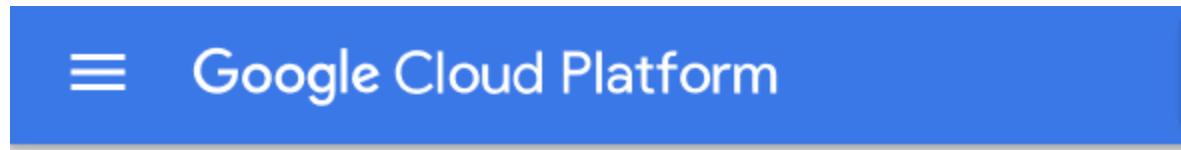


Image courtesy <https://console.cloud.google.com/home>

Give your app a name and make a note of your project ID.



New Project



You have 24 projects remaining in your quota. [Learn more.](#)

Project name [?](#)

myHelloWorld

Your project ID will be myhelloworld-201222 [?](#) [Edit](#)

[Create](#)

[Cancel](#)

Image courtesy.

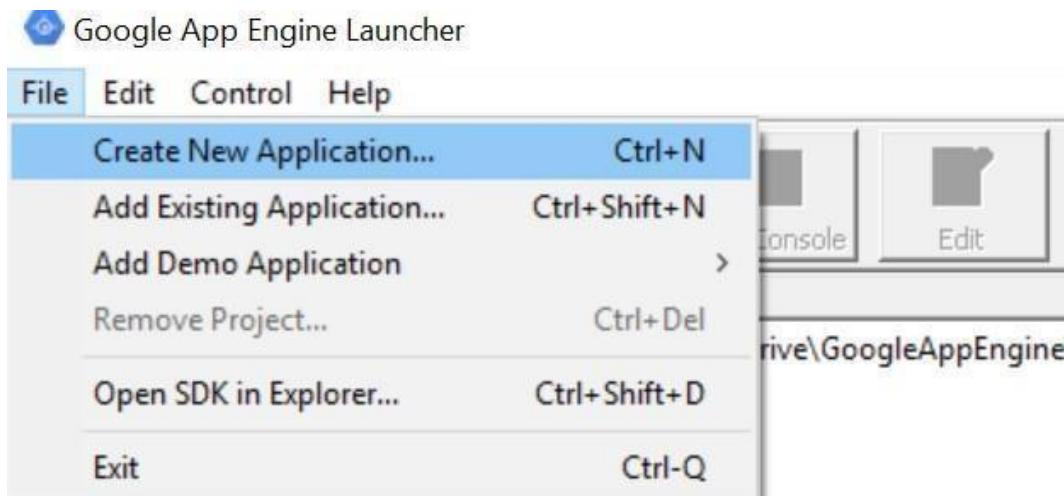
<https://console.cloud.google.com/home>

Hit the create button and Google should take a few minutes to set up all that is necessary for your newly created app.

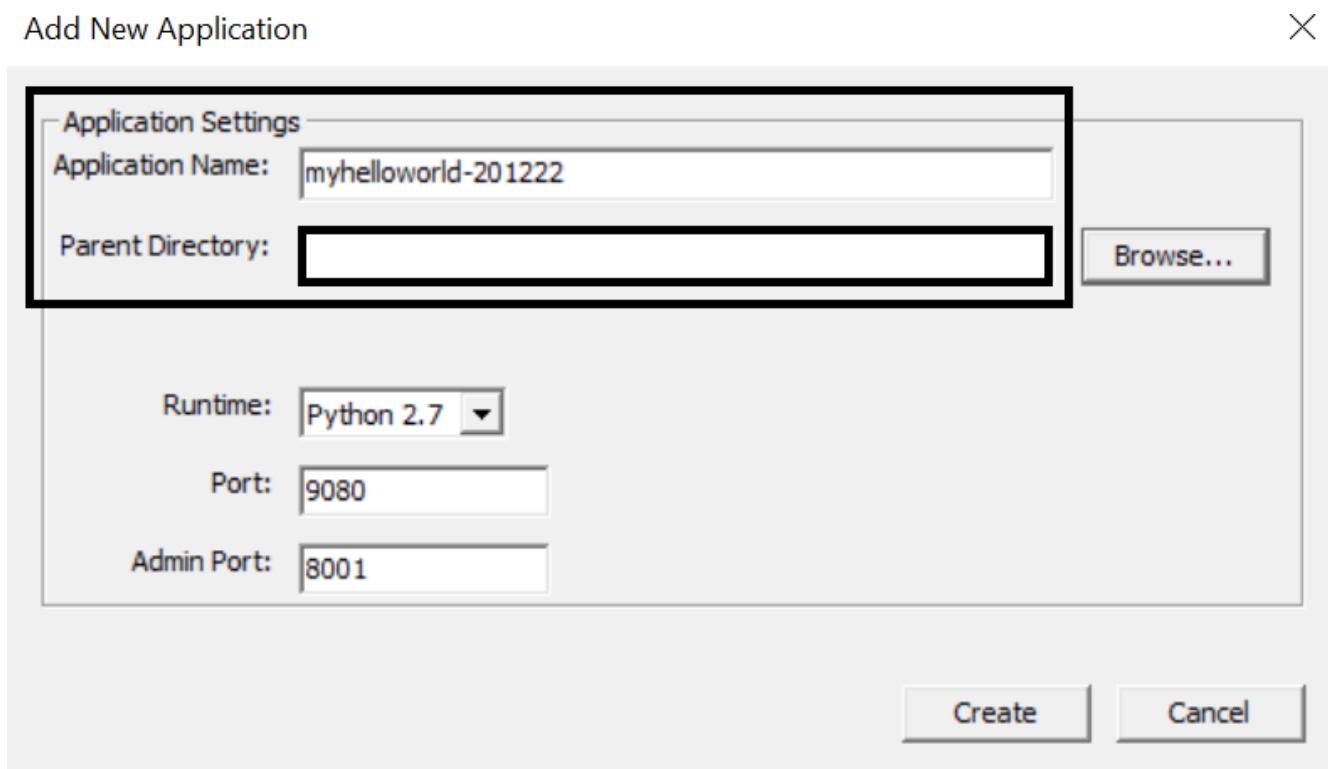
Step 4. Fork the app to develop it locally

The next step in the process is to fork the app on your local machine. This will allow you to make changes to the app locally and deploy it whenever you wish to.

Go to Google App Engine launcher and create a new application.



Enter the project ID of your newly created app. Also, provide the folder (local destination) where you wish to store the app locally. Make sure you select the Python 2.7 as your runtime engine.

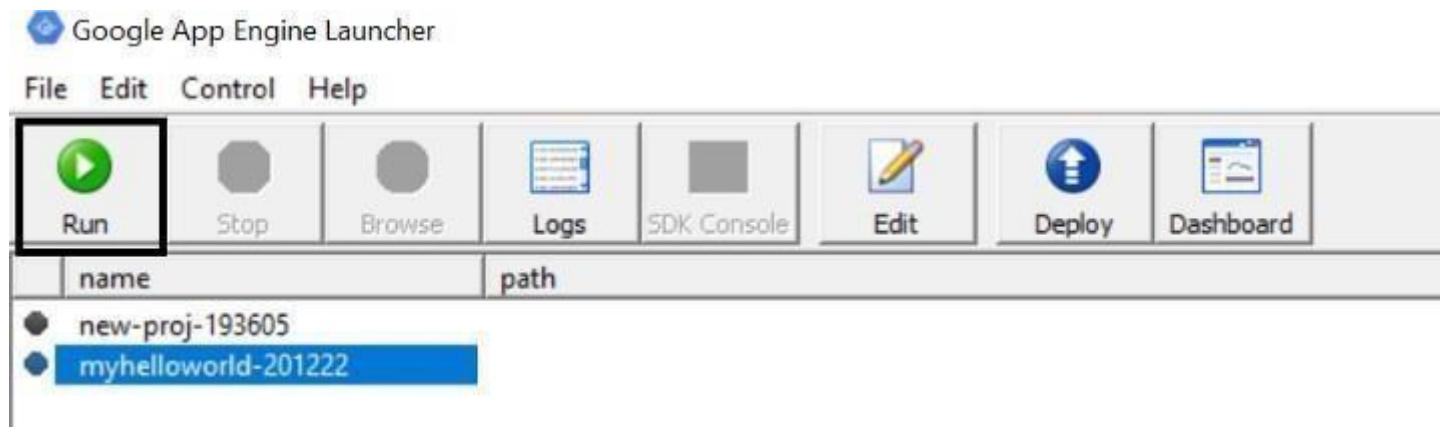


Hit the create button, and you should see your app listed on the window that follows. You should also check that you now see some files in your local storage (the directory you chose in the screenshot above) after this step.

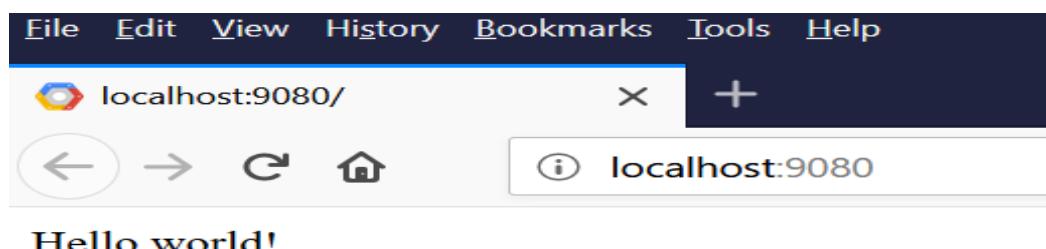
Step 5. Run the app locally

Before you go ahead and make some changes to the app, it is important to check whether or not you have executed all the above steps correctly. This can be done by simply running the app locally.

Select the app and hit the run button on the window.



Wait for a few seconds until you can hit the **Browse** button. Once the **Browse** button becomes clickable, click it. This should take you to the browser, and you should see the hello world text appear in your browser window. Alternatively, you can manually go to the browser and use the port specified to access the app.



As long as you see the above screen, you are all set.

Step 6. Understand the app structure

It is finally time to look at the lines of code which are running this webapp. Open your app folder in the text editor of your choice. I recommend [Sublime text](#) or [VS Code](#). However, feel free to choose the one you prefer.

Here is a description of the various files.

app.yaml

This file is a basic markup file that stores information (some metadata) about the app. It is important to note the following crucial parts of the file.

1. application

This is the project ID which you should never change. This is the unique identifier for the app

2. url -> script

This is the homepage for the app. In other words, this file will be rendered in your browser when you launch the app

3. libraries

This is where you can include external libraries to use within the webapp

```
application: myhelloworld-201222
version: 1
runtime: python27
api_version: 1
threadsafe: yes

handlers:
- url: /favicon\.ico
  static_files: favicon.ico
  upload: favicon\.ico

- url: .*
  script: main.app

libraries:
- name: webapp2
  version: "2.5.2"
```

file in the webapp folder

main.py program

This is the homepage of the app (as discussed above). Note that the hello world text in the browser window (step 5) is due to the code you see highlighted below.

The screenshot shows a code editor interface with a dark theme. On the left, there's a sidebar titled "FOLDERS" containing project structures like "new-proj-193605", "scratchpad", "hw2", and "myhelloworld-201222". Inside "myhelloworld-201222", files "app.yaml", "favicon.ico", "index.yaml", and "main.py" are listed. The "main.py" file is currently open in the main editor area. The code is as follows:

```
1  #!/usr/bin/env python
2  #
3  # Copyright 2007 Google Inc.
4  #
5  # Licensed under the Apache License, Version 2.0 (the "License");
6  # you may not use this file except in compliance with the License.
7  # You may obtain a copy of the License at
8  #
9  #     http://www.apache.org/licenses/LICENSE-2.0
10 #
11 # Unless required by applicable law or agreed to in writing, software
12 # distributed under the License is distributed on an "AS IS" BASIS,
13 # WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
14 # See the License for the specific language governing permissions and
15 # limitations under the License.
16 #
17 import webapp2
18
19 class MainHandler(webapp2.RequestHandler):
20     def get(self):
21         self.response.write('Hello world!')
22
23 app = webapp2.WSGIApplication([
24     ('/', MainHandler)
25 ], debug=True)
26
```

main.py file in the webapp folder

Step 7. Make your changes and deploy the new app

No hello world app is ever complete without the developer changing the hello world text to something else just to make sure that everything happening behind the scenes is working as it should.

Go ahead and change the text in the above screenshot to something else.

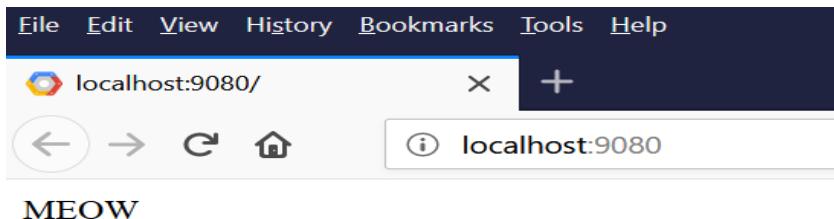
```

17 import webapp2
18
19 class MainHandler(webapp2.RequestHandler):
20     def get(self):
21         self.response.write('MEOW')
22
23 app = webapp2.WSGIApplication([
24     ('/', MainHandler)
25 ], debug=True)
26

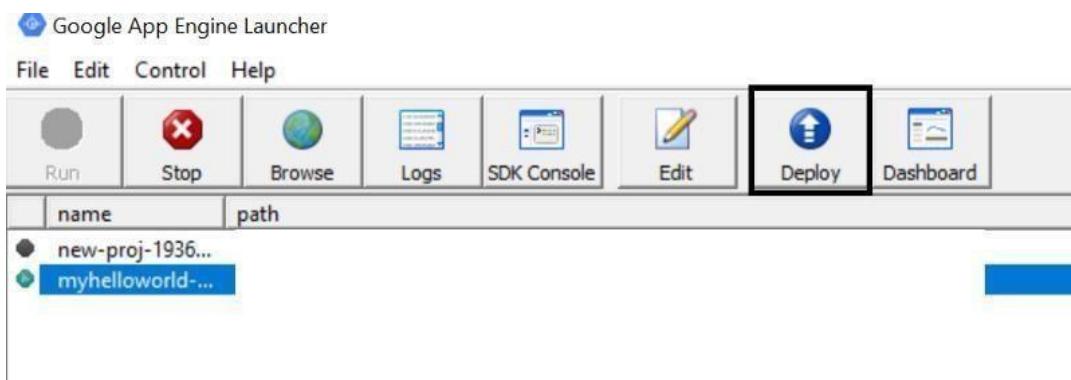
```

Output:

Save the changes, go to the browser and refresh the page. You should see the page with the text “MEOW” displayed.



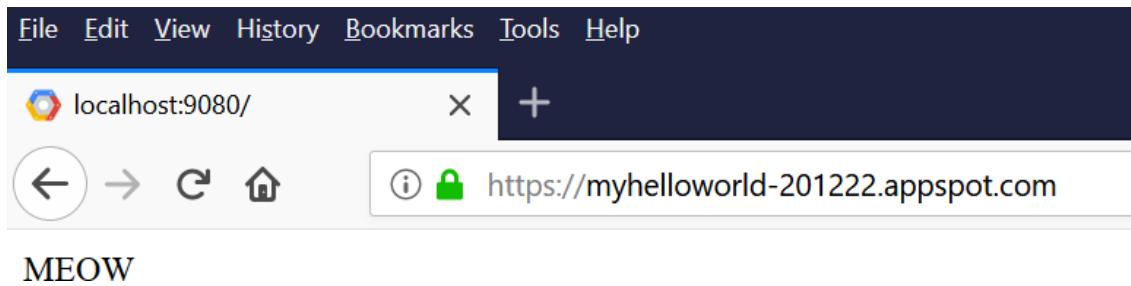
Finally, it is time to deploy your changes to the cloud to make them globally accessible via a URL. Go to the App Engine launcher, select the app, and hit the **Deploy** button.



This will ensure your app gets deployed onto Google Cloud. To check whether or not everything worked just fine, go to the URL below:

<https://<yourProjectID>.appspot.com/>

You should see the exact same window as above, except now, it is a URL that is globally accessible.



Result:

EXP NO: 5
DATE:

**SIMULATE A CLOUD SCENARIO USING CLOUD SIM AND RUN A
SCHEDULING ALGORITHM THAT IS NOT PRESENT IN CLOUD SIM**

Aim:

Procedure:

CloudSim is written in Java. The knowledge you need to use CloudSim is basic Java programming and some basics about cloud computing. Knowledge of programming IDEs such as Eclipse or NetBeans is also helpful. It is a library and, hence, CloudSim does not have to be installed. Normally, you can unpack the downloaded package in any directory, add it to the Java classpath and it is ready to be used. Please verify whether Java is available on your system.

To use CloudSim in Eclipse:

1. Download CloudSim installable files from <https://code.google.com/p/cloudsim/downloads/list> and unzip
2. Open Eclipse
3. Create a new Java Project: File -> New
4. Import an unpacked CloudSim project into the new Java Project
5. The first step is to initialise the CloudSim package by initialising the CloudSim library, as follows:

```
CloudSim.init(num_user, calendar, trace_flag)
```

6. Data centres are the resource providers in CloudSim; hence, creation of data centres is a second step. To create Datacenter, you need the DatacenterCharacteristics object that stores the properties of a data centre such as architecture, OS, list of machines, allocation policy that covers the time or space shared, the time zone and its price:

```
Datacenter datacenter9883 = new Datacenter(name, characteristics,  
new VmAllocationPolicySimple(hostList),s
```

7. The third step is to create a broker:

```
DatacenterBroker broker = createBroker();
```

8. The fourth step is to create one virtual machine unique ID of the VM, userId ID of the VM's owner, mips, number Of Pes amount of CPUs, amount of RAM, amount of bandwidth, amount of storage, virtual machine monitor, and cloudletScheduler policy for cloudlets:

```
Vm vm = new Vm(vmid, brokerId, mips, pesNumber, ram, bw, size, vmm, new  
CloudletSchedulerTimeShared())
```

9. Submit the VM list to the broker:

```
broker.submitVmList(vmlist)
```

10. Create a cloudlet with length, file size, output size, and utilisationmodel:

```
Cloudlet cloudlet = new Cloudlet(id, length, pesNumber, fileSize, outputSize, utilizationModel,  
utilizationMode
```

11. Submit the cloudlet list to the broker:

```
broker.submitCloudletList(cloudletList)
```

12. Start

the simulation:

```
CloudSim.startSimulation()
```

Sample Output from the Existing Example:

Starting

CloudSimExample1...

Initialising...

Starting CloudSim version 3.0 Datacenter_0 is starting...

Broker is

starting...

Entities

started.

: Broker: Cloud Resource List received with

1resource(s) 0.0: Broker: Trying to Create VM #0

inDatacenter_0

: Broker: VM #0 has been created in Datacenter #2, Host#0

0.1: Broker: Sending cloudlet 0 to VM #0

400.1: Broker: Cloudlet 0 received

: Broker: All Cloudlets

executed.Finishing.... 400.1: Broker:

Destroying VM#0

Broker is shutting down...

Simulation: No more

future events

CloudInformationService: Notify all CloudSim entities

shutting down. D

Broker

down...

Sim

ed.

OUTPUT:

Cloudlet ID	STATUS	Data center ID	VM ID	Time	Start Time	FinishTime
0	SUCCESS	20	400	0.1		400.1
*****Datacenter:Datacenter_0*****						
		Userid		Debt		
3			35.6			

CloudSimExample1 finished!

RESULT:

EXP NO: 6 DATE:	FIND A PROCEDURE TO TRANSFER THE FILES FROM ONE VIRTUAL MACHINE TO ANOTHER VIRTUAL MACHINE
----------------------------------	---

Aim:**Procedure:****Method1: Creating a Shared Folder in VirtualBox**

A shared folder is a folder which makes its files available on both the guest machine *and* the host machine **at the same time**. Creating a shared folder between the guest and the host allows you to easily manage files which should be present on both machines. The course virtual machines are ready to use sharedfolders right away, but if you are using the virtual machine on your personal computer you will need to specify which folder to use as shared storage.

Shared Folders on SCS Lab Computers using Course VMs:

If you are using a course VM on a lab computer, it is likely that a shared folder has already been setup for you. On the desktop of your course VM you should notice a folder titled *SharedFolders*. Inside of this you will find any folders that have been shared between the course VM and lab computers.

You should see two folders that have already been configured for you: **Z_DRIVE** and **Temp**.

Z_DRIVE gives you access to your [Windows Account](#) `z:\ drive`. This is storage that is persistent to your SCS account and available as a network drive on the lab computers.

Temp gives you access to the folder found at `D:\temp` on the lab computer. Files stored in this folder are local to the machine, meaning that they can be accessed **faster**, but will **delete** from the system when you log out.

If you are working with data that you will need to use again, use the *Z_DRIVE* for your shared folder. If you need faster read/write speed, use the *Temp* folder, but remember to backup your files or they will be deleted when you log off the computer.

Shared Folders on Personal Computers:

If you are using your own personal machine, you will need to configure VirtualBox to look in the right place for your shared files.

First, click on the guest machine you intend to share files with. From there, you can select the *guestSettings* and navigate to *Shared Folders* on the left side menu. To create a new shared folder, either click the *New Folder* icon on the right menu **or** right click the empty list of shared folders and click *Add Shared Folder*. From here, there are six options:

- **Folder Path:** The folder name on the **host** machine. Click the drop down menu and navigate to the folder you would like to share.
- **Folder Name:** This is the name of the folder as it will appear on the **guest** machine.
- **Read-Only:** If you check read-only, the **guest** machine will be unable to write changes to the folder. This is valuable when you only want to send files *to* the virtual machine, but do not want to risk having the files modified by the guest.
- **Auto-Mount:** When any external storage is connected to a computer it must be *mounted* in order to be used. It is recommended that you turn on auto-mounting, unless you are familiar with the process of mounting a drive yourself.
- **Mount Point:** Unless you already know about mount points, leave this blank.
- **Make Permanent:** If you check this, the shared folder will be a permanent **machine folder**. If it is not checked, the folder will not be shared after a shutdown.

On the course virtual machines, when you load into the desktop, you should see a folder labelled *SharedFolders*. In there you will see any folders that are currently mounted and being shared.

Steps:

1. Select the guest machine you wish to share files with

2. Click Settings > Shared Folders
3. Right-click and select Add Shared Folder and use the following settings:

Folder Path: Click the dropdown arrow, select **Other**, and navigate to the folder you would like to share

Folder Name: Anything to identify it on the guest machine

Read-Only: Unchecked (Checked, if you are exclusively pulling files **from the host**)

Auto-Mount: Checked

Mount Point: Leave blank

4. Click OK

Method 2: Dragging and Dropping Files in VirtualBox

If you only need to transfer a few files quickly, you can simply drag and drop the files in. On the top bar of the running guest machine, click on *Devices* > *Drag and Drop* and make sure that *Bidirectional* is selected. This means that you will be able to drag files from the host to the guest and from the guest to the host. Once bidirectional drag and drop is checked, you should be able to begin dragging and dropping files.

NOTE: Sometimes when dragging files **into the course VM**, you may not be able to drag into the file browser directly. If you encounter this issue, you should drag your files onto the *Desktop* and move the files around from there. You should see the cursor change when it is ready to drop files.

You can also drag files from the guest machine into the host. To do this, simply open the file browser on the host to where you would like to drop the files and drag the files from the virtual machine into the file browser of the host. File transfers should be pretty quick; if the virtual machine seems stuck when transferring, simply cancel the transfer and try again.

Method 3: Managing Files with NextCloud

On any virtual machine, including VirtualBox, VMWare, or the virtual machines hosted on the [SCS OpenStack](#), you can access the [SCS NextCloud](#) services to move files between multiple machines and your [SCS Windows Account storage](#). NextCloud offers you all of your SCS storage in one remote location, similar to how you might use other file hosting services like Dropbox or Google Drive. Before trying to use NextCloud, you should check that you can access the service by [logging in here](#).

If you can access the NextCloud services, you can browse the various file storage services available to you:

- **Linux Home:** These are the files from your [SCS Linux Account](#)
- **Windows Home:** These are the files from your [SCS Windows Account](#) and your lab `z:\` drive.
- **NextCloud:** In addition to the other storage accounts provided to you by the SCS, you can also upload up to 20GB of files directly to NextCloud.

With NextCloud, you can upload your files from any machine with an internet connection and download them onto any other machine with an internet connection. For example, you can move project files off of your virtual machine, onto the NextCloud storage, and then download them on your personal laptop.

Alternatively, you can upload files from your personal PC onto the NextCloud storage, place it into the *Windows Home* folder, and access those files from either the lab `z:\` drive or download them on a virtual machine like VirtualBox or OpenStack.

Unloading Files to NextCloud from a Lab Computer:

If you would like to upload files from a lab computer, the easiest way to do this is to place the files you would like to transfer into your `z:\` drive. These files will be automatically backup into your NextCloud storage under the *Windows Home* folder. After that, you can move them into the main NextCloud storage or choose to keep them in your `z:\` drive.

Uploading Files to NextCloud from a VM or Other PC:

If you would like to upload files from either a VM or any other computer, you can login to the NextCloud service using any of the available interfaces, such as the [web interface](#). Press the “+” icon in the top left of the file browser and select *Upload File*. From here, you can choose to keep it in the main NextCloud storage, move it into your Windows Account storage (the *Windows Home* folder), or into your Linux Account storage (the *Linux Home* folder).

Downloading NextCloud Files to a VM or Other PC:

Once your files are uploaded you will be able to download those files onto any machine which can connect to NextCloud. First, log in to your preferred NextCloud interface (eg. the [web interface](#)). Navigate to the folder which contains the files you would like to download. Once you are in the target folder, click the checkbox next to each file you would like to download. Above the file listing you should notice the context bar

changing to tell you how many files you have selected and a button labelled *Actions*. Click *Actions > Download*.

If you have selected a single file, it will prompt you to confirm the download. If you have chosen more than one file, NextCloud will place all of the selected files into a *zip archive*. Before you can use the files, you will need to extract them from the archive. Once you have downloaded your file, or extracted your archive, you are ready to use your files on your machine.

Result:

Exp No: 7(A)
DATE:

INSTALL HADOOP SINGLE NODE CLUSTER AND RUN SIMPLE APPLICATIONS LIKE WORD COUNT

Aim:

Step 1: To download the Java 8 Package. Save this file in your home directory.

Step 2: Extract the Java Tar File.

Command: tar -xvf jdk-8u101-linux-i586.tar.gz

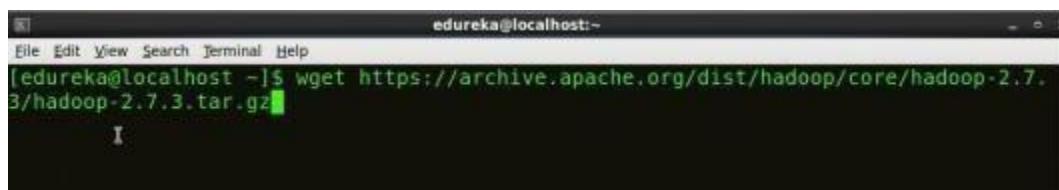


```
edureka@localhost:~$ tar -xvf jdk-8u101-linux-i586.tar.gz
```

Fig: Hadoop Installation – Extracting Java Files

Step 3: Download the Hadoop 2.7.3 Package.

Command: wget <https://archive.apache.org/dist/hadoop/core/hadoop-2.7.3/hadoop-2.7.3.tar.gz>

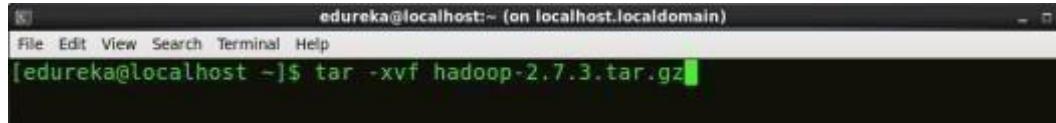


```
edureka@localhost:~$ wget https://archive.apache.org/dist/hadoop/core/hadoop-2.7.3/hadoop-2.7.3.tar.gz
```

Fig: Hadoop Installation – Downloading Hadoop

Step 4: Extract the Hadoop tar File.

Command: tar -xvf hadoop-2.7.3.tar.gz



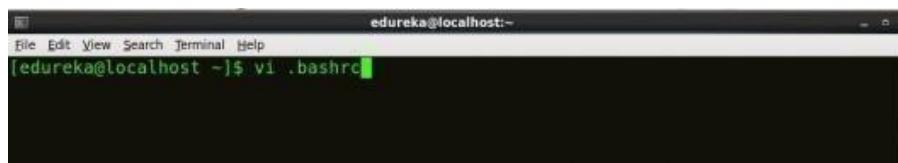
```
edureka@localhost:~ (on localhost.localdomain)
File Edit View Search Terminal Help
[edureka@localhost ~]$ tar -xvf hadoop-2.7.3.tar.gz
```

Fig: Hadoop Installation – Extracting Hadoop Files

Step 5: Add the Hadoop and Java paths in the bash file (.bashrc). Open.

bashrc file. Now, add Hadoop and Java Path as shown below.

Command: vi .bashrc



```
# User specific aliases and functions

export HADOOP_HOME=$HOME/hadoop-2.7.3
export HADOOP_CONF_DIR=$HOME/hadoop-2.7.3/etc/hadoop
export HADOOP_MAPRED_HOME=$HOME/hadoop-2.7.3
export HADOOP_COMMON_HOME=$HOME/hadoop-2.7.3
export HADOOP_HDFS_HOME=$HOME/hadoop-2.7.3
export YARN_HOME=$HOME/hadoop-2.7.3
export PATH=$PATH:$HOME/hadoop-2.7.3/bin

# Set JAVA_HOME

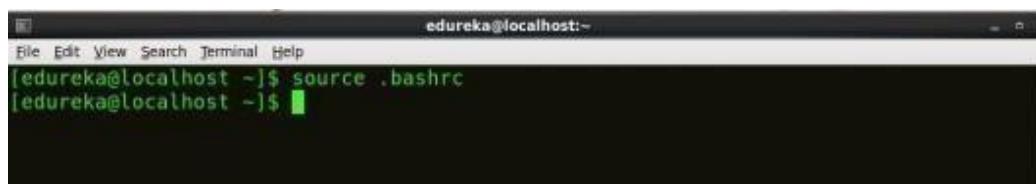
export JAVA_HOME=/home/edureka/jdk1.8.0_101
export PATH=/home/edureka/jdk1.8.0_101/bin:$PATH
```

Fig: Hadoop Installation – Setting Environment Variable

Then, save the bash file and close it.

For applying all these changes to the current Terminal, execute the source command.

Command: source .bashrc

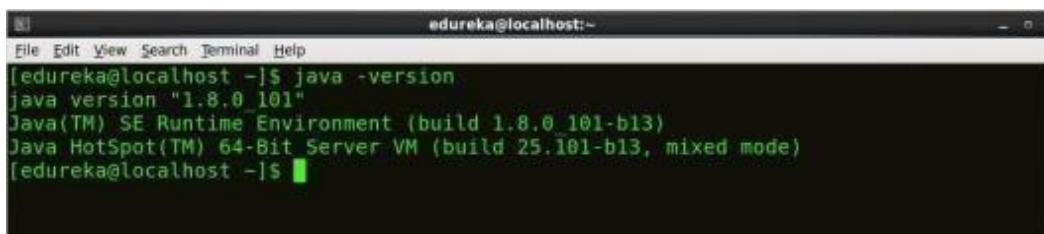


```
edureka@localhost:~
File Edit View Search Terminal Help
[edureka@localhost ~]$ source .bashrc
[edureka@localhost ~]$
```

Fig: Hadoop Installation – Refreshing environment variables

To make sure that Java and Hadoop have been properly installed on your system and can be accessed through the Terminal, execute the java -version and hadoop version commands.

Command: java -version

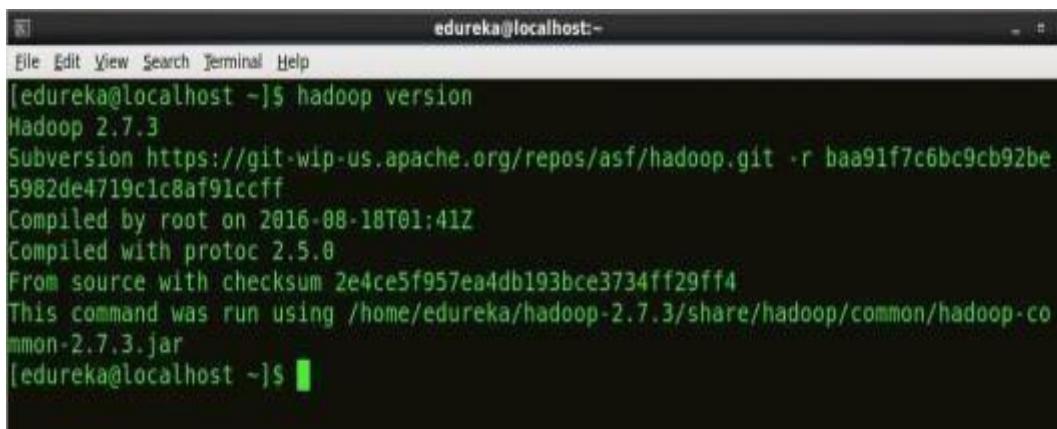


```
edureka@localhost:~$ java -version
java version "1.8.0_101"
Java(TM) SE Runtime Environment (build 1.8.0_101-b13)
Java HotSpot(TM) 64-Bit Server VM (build 25.101-b13, mixed mode)
[edureka@localhost ~]$
```

A screenshot of a terminal window titled 'edureka@localhost:~'. The window shows the command 'java -version' being run and its output. The output indicates Java version 1.8.0_101, Java(TM) SE Runtime Environment (build 1.8.0_101-b13), and Java HotSpot(TM) 64-Bit Server VM (build 25.101-b13, mixed mode). The terminal prompt '[edureka@localhost ~]\$' is visible at the bottom.

Fig: Hadoop Installation – Checking Java Version

Command: hadoop version



```
edureka@localhost:~$ hadoop version
Hadoop 2.7.3
Subversion https://git-wip-us.apache.org/repos/asf/hadoop.git -r baa91f7c6bc9cb92be
5982de4719c1c8af91ccff
Compiled by root on 2016-08-18T01:41Z
Compiled with protoc 2.5.0
From source with checksum 2e4ce5f957ea4db193bce3734ff29ff4
This command was run using /home/edureka/hadoop-2.7.3/share/hadoop/common/hadoop-co
mmon-2.7.3.jar
[edureka@localhost ~]$
```

A screenshot of a terminal window titled 'edureka@localhost:~'. The window shows the command 'hadoop version' being run and its output. The output displays Hadoop version 2.7.3, Subversion details, compilation date (2016-08-18T01:41Z), and the checksum of the source code. The terminal prompt '[edureka@localhost ~]\$' is visible at the bottom.

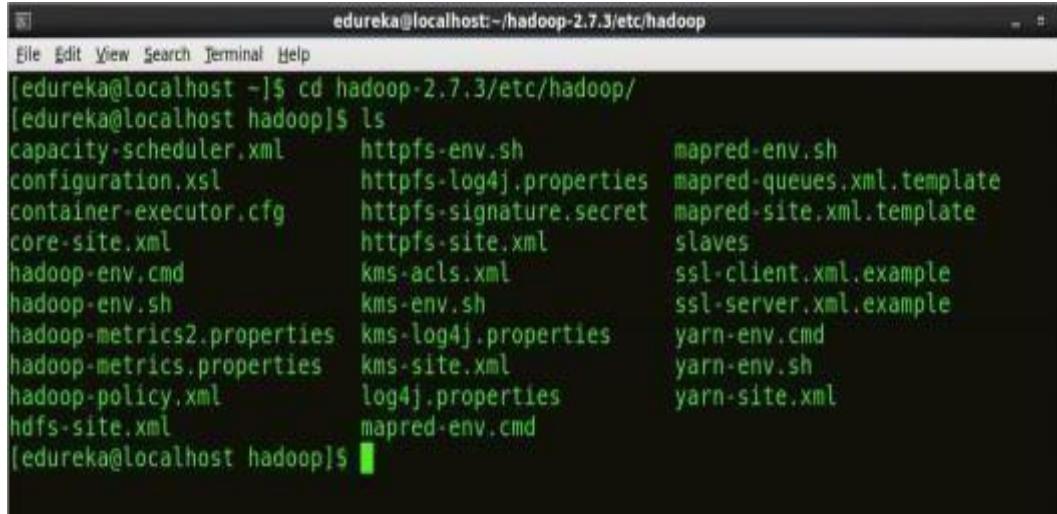
Fig: Hadoop Installation – Checking Hadoop Version

Step 6: Edit the **Hadoop Configuration files**.

Command: cd hadoop-2.7.3/etc/hadoop/

Command: ls

All the Hadoop configuration files are located in **hadoop-2.7.3/etc/hadoop** directory as you can see in the snapshot below:



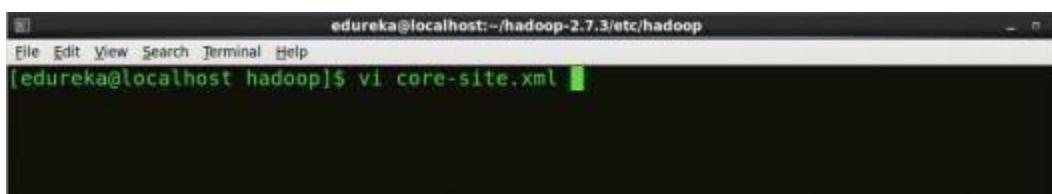
```
edureka@localhost:~/hadoop-2.7.3/etc/hadoop
[edureka@localhost ~]$ cd hadoop-2.7.3/etc/hadoop/
[edureka@localhost hadoop]$ ls
capacity-scheduler.xml      httpfs-env.sh          mapred-env.sh
configuration.xsl            httpfs-log4j.properties  mapred-queues.xml.template
container-executor.cfg       httpfs-signature.secret mapred-site.xml.template
core-site.xml                httpfs-site.xml        slaves
hadoop-env.cmd               kms-acls.xml         ssl-client.xml.example
hadoop-env.sh                kms-env.sh           ssl-server.xml.example
hadoop-metrics2.properties   kms-log4j.properties  yarn-env.cmd
hadoop-metrics.properties    kms-site.xml         yarn-env.sh
hadoop-policy.xml            log4j.properties     yarn-site.xml
hdfs-site.xml                mapred-env.cmd
```

Fig: Hadoop Installation – Hadoop Configuration Files

Step 7: Open *core-site.xml* and edit the property mentioned below inside configuration tag:

core-site.xml informs Hadoop daemon where NameNode runs in the cluster. It contains configuration settings of Hadoop core such as I/O settings that are common to HDFS & MapReduce.

Command: vi core-site.xml



```
edureka@localhost:~/hadoop-2.7.3/etc/hadoop
[edureka@localhost hadoop]$ vi core-site.xml
```



```
<configuration>
<property>
<name>fs.default.name</name>
<value>hdfs://localhost:9000</value>
</property>
</configuration>
```

Fig: Hadoop Installation – Configuring core-site.xml

```
1           <?xmlversion="1.0"encoding="UTF-8"?>
2   <?xmlstylesheet type="text/xsl" href="configuration.xsl"?>
3       <configuration>
4           <property>
5               <name>fs.default.name</name>
6               <value>hdfs://localhost:9000</value>
7           </property>
8       </configuration>
```

Step 8: Edit *hdfs-site.xml* and add the property mentioned below inside configuration tag:

```
<configuration>
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
<property>
<name>dfs.permission</name>
<value>false</value>
</property>
```

hdfs-site.xml contains configuration settings of HDFS daemons (i.e. NameNode, DataNode, Secondary NameNode). It also includes the replication factor and block size of HDFS.

Command: vi hdfs-site.xml

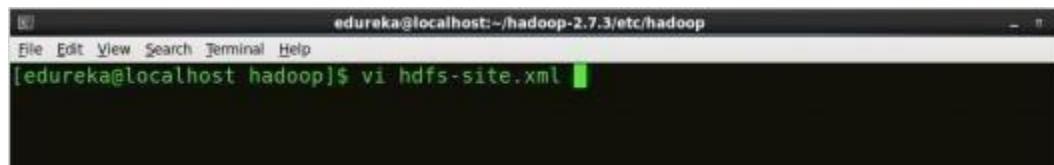


Fig: Hadoop Installation – Configuring hdfs-site.xml

```
1           <?xmlversion="1.0"encoding="UTF-8"?>
2   <?xml-stylesheettype="text/xsl"href="configuration.xsl"?>
3       <configuration>
4           <property>
5               <name>dfs.replication</name>
6               <value>1</value>
7           </property>
8           <property>
9               <name>dfs.permission</name>
```

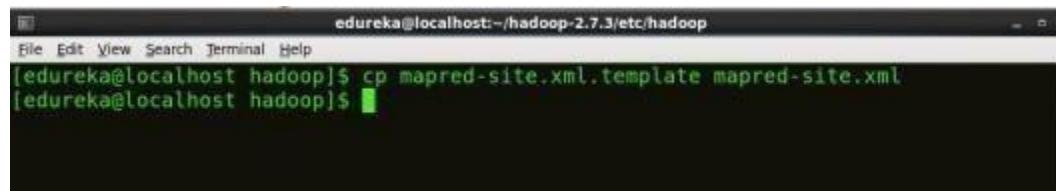
Step 9: Edit the *mapred-site.xml* file and edit the property mentioned below inside configuration tag:

mapred-site.xml contains configuration settings of MapReduce application like number of JVM that can run in parallel, the size of the mapper and the reducer process, CPU cores available for a process,etc.

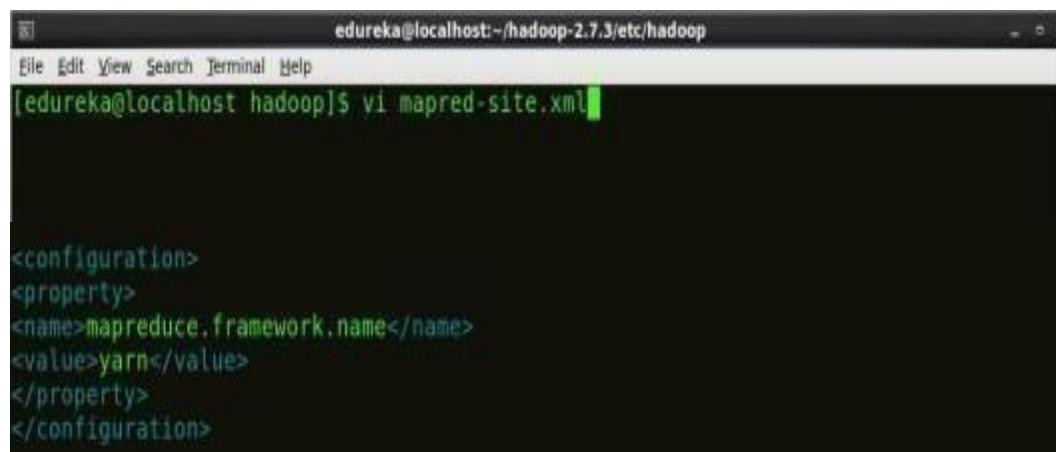
In some cases, *mapred-site.xml* file is not available. So, we have to create the *mapred-site.xml* file using *mapred-site.xml* template.

Command: cp *mapred-site.xml.template* *mapred-site.xml*

Command: vi *mapred-site.xml*.



```
edureka@localhost:~/hadoop-2.7.3/etc/hadoop
File Edit View Search Terminal Help
[edureka@localhost hadoop]$ cp mapred-site.xml.template mapred-site.xml
[edureka@localhost hadoop]$
```



```
edureka@localhost:~/hadoop-2.7.3/etc/hadoop
File Edit View Search Terminal Help
[edureka@localhost hadoop]$ vi mapred-site.xml
```



```
<configuration>
<property>
<name>mapreduce.framework.name</name>
<value>yarn</value>
</property>
</configuration>
```

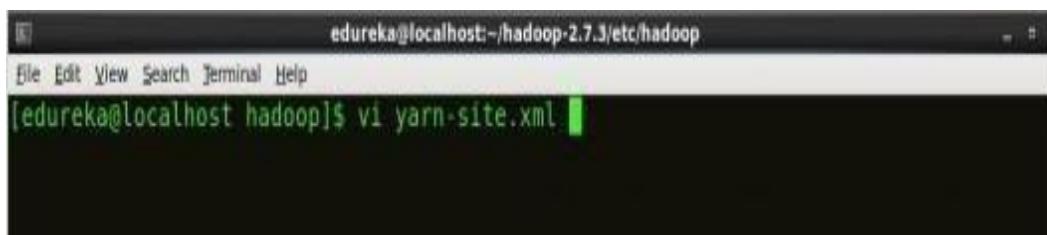
Fig: Hadoop Installation – Configuring mapred-site.xml

```
1             <?xmlversion="1.0"encoding="UTF-8"?>
2     <?xmlstylesheettype="text/xsl"href="configuration.xsl"?>
3         <configuration>
4             <property>
5                 <name>mapreduce.framework.name</name>
6                     <value>yarn</value>
7                         </property>
8                     </configuration>
```

Step 10: Edit *yarn-site.xml* and edit the property mentioned below inside configuration tag:

yarn-site.xml contains configuration settings of ResourceManager and NodeManager like application memory management size, the operation needed on program & algorithm, etc.

Command: vi *yarn-site.xml*



```
<configuration>
<property>
<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
</property>
<property>
<name>yarn.nodemanager.auxservices.mapreduce.shuffle.class</name>
<value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
</configuration>
```

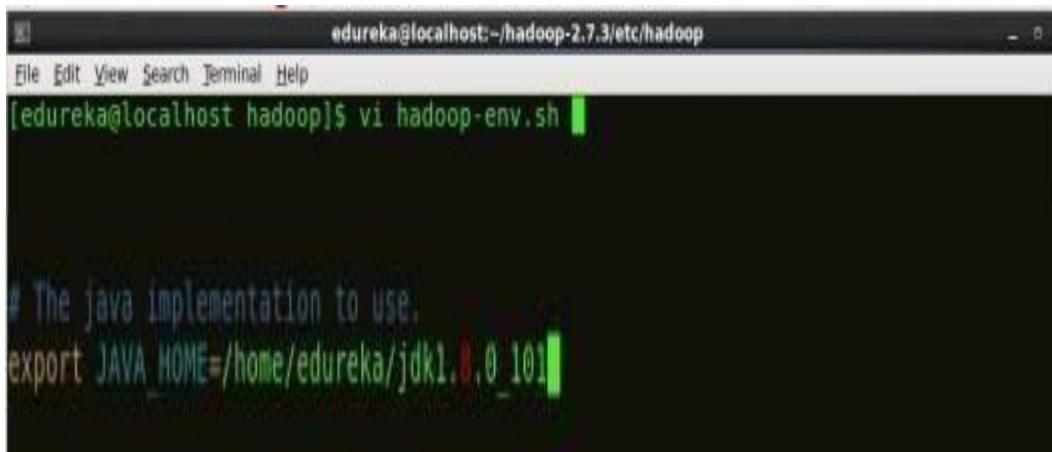
*Fig: Hadoop Installation – Configuring *yarn-site.xml**

```
1
2
3         <?xmlversion="1.0">
4             <configuration>
5                 <property>
6                     <name>yarn.nodemanager.aux-services</name>
7                         <value>mapreduce_shuffle</value>
8                     </property>
9                     <property>
10                         <name>yarn.nodemanager.auxservices.mapreduce.shuffle.class</name>
11                         <value>org.apache.hadoop.mapred.ShuffleHandler</value>
12                     </property>
13             </configuration>
14
```

Step 11: Edit *hadoop-env.sh* and add the Java Path as mentioned below:

hadoop-env.sh contains the environment variables that are used in the script to run Hadoop like Java home path, etc.

Command: vi *hadoop-env.sh*



```
[edureka@localhost hadoop]$ vi hadoop-env.sh
# The java implementation to use.
export JAVA_HOME=/home/edureka/jdk1.8.0_101
```

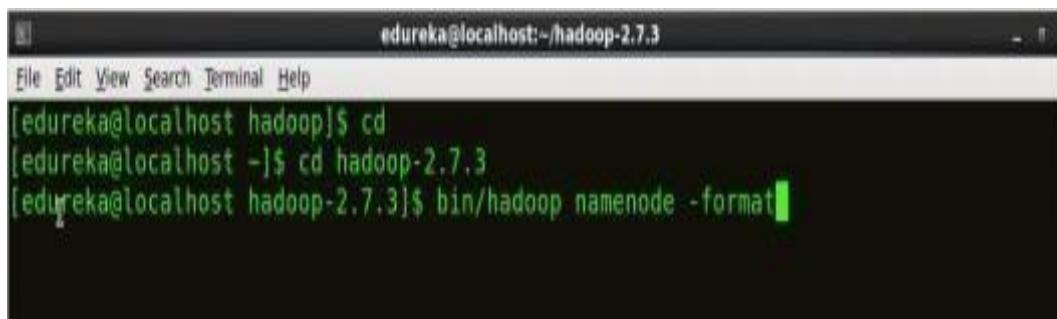
Fig: Hadoop Installation – Configuring hadoop-env.sh

Step 12: Go to Hadoop home directory and format the NameNode.

Command: cd

Command: cd hadoop-2.7.3

Command: bin/hadoop namenode -format



The screenshot shows a terminal window titled 'edureka@localhost:~/hadoop-2.7.3'. The window has a menu bar with 'File', 'Edit', 'View', 'Search', 'Terminal', and 'Help'. The main area of the terminal shows the following command sequence:
[edureka@localhost hadoop]\$ cd
[edureka@localhost ~]\$ cd hadoop-2.7.3
[edureka@localhost hadoop-2.7.3]\$ bin/hadoop namenode -format

Fig: Hadoop Installation – Formatting NameNode

This formats the HDFS via NameNode. This command is only executed for the first time. Formatting the file system means initializing the directory specified by the `dfs.name.dir` variable.

Never format, up and running Hadoop filesystem. You will lose all your data stored in theHDFS.

Step 13: Once the NameNode is formatted, go to `hadoop-2.7.3/sbin` directory and start all the daemons.

Command: cd hadoop-2.7.3/sbin

Either you can start all daemons with a single command or do it individually.

Command: ./start-all.sh

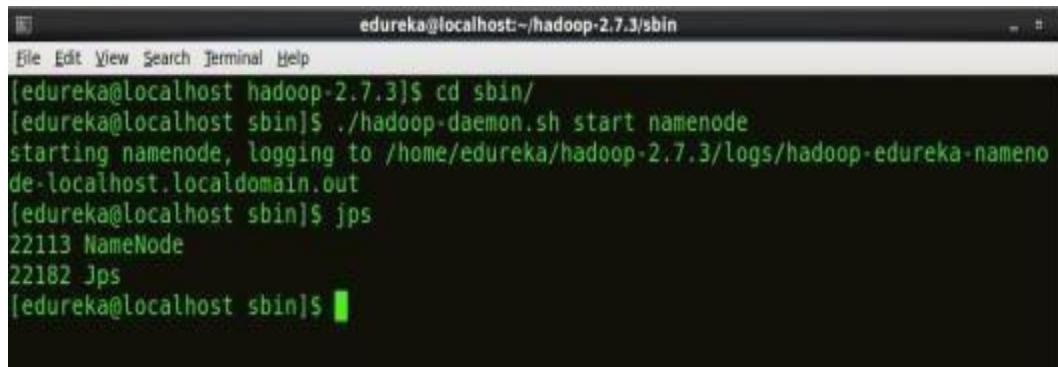
The above command is a combination of `start-dfs.sh`, `start-yarn.sh` & `mr-jobhistory-daemon.sh`

Or you can run all the services individually as below:

Start NameNode:

The NameNode is the centerpiece of an HDFS file system. It keeps the directory tree of all files stored in the HDFS and tracks all the file stored across the cluster.

Command: ./hadoop-daemon.sh start namenode



A screenshot of a terminal window titled "edureka@localhost:~/hadoop-2.7.3/sbin". The window shows the command ./hadoop-daemon.sh start namenode being run, followed by the output indicating the NameNode is starting and logging to a specific log file. The "jps" command is then run to show the process IDs for the NameNode and JPS.

```
[edureka@localhost hadoop-2.7.3]$ cd sbin/  
[edureka@localhost sbin]$ ./hadoop-daemon.sh start namenode  
starting namenode, logging to /home/edureka/hadoop-2.7.3/logs/hadoop-edureka-namenode-localhost.localdomain.out  
[edureka@localhost sbin]$ jps  
22113 NameNode  
22182 Jps  
[edureka@localhost sbin]$
```

StartDataNode:

– Starting NameNode

On startup, a DataNode connects to the Namenode and it responds to the requests from the Namenode for different operations.

Command: ./hadoop-daemon.sh start datanode



A screenshot of a terminal window titled "edureka@localhost:~/hadoop-2.7.3/sbin". The window shows the command ./hadoop-daemon.sh start datanode being run, followed by the output indicating the DataNode is starting and logging to a specific log file. The "jps" command is then run to show the process IDs for the NameNode, JPS, and DataNode.

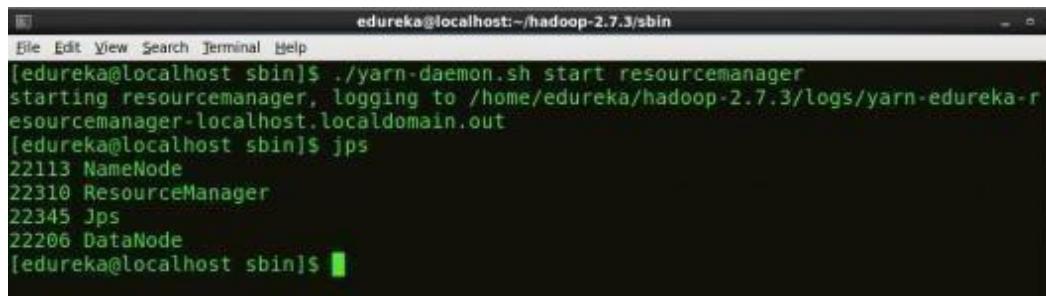
```
[edureka@localhost sbin]$ ./hadoop-daemon.sh start datanode  
starting datanode, logging to /home/edureka/hadoop-2.7.3/logs/hadoop-edureka-datanode-localhost.localdomain.out  
[edureka@localhost sbin]$ jps  
22113 NameNode  
22278 Jps  
22206 DataNode  
[edureka@localhost sbin]$
```

Fig: Hadoop Installation – Starting DataNode

Start ResourceManager:

ResourceManager is the master that arbitrates all the available cluster resources and thus helps in managing the distributed applications running on the YARN system. Its work is to manage each NodeManagers and the each application's ApplicationMaster.

Command: ./yarn-daemon.sh start resourcemanager



```
edureka@localhost:~/hadoop-2.7.3/sbin
File Edit View Search Terminal Help
[edureka@localhost sbin]$ ./yarn-daemon.sh start resourcemanager
starting resourcemanager, logging to /home/edureka/hadoop-2.7.3/logs/yarn-edureka-resourcemanager-localhost.localdomain.out
[edureka@localhost sbin]$ jps
22113 NameNode
22310 ResourceManager
22345 Jps
22206 DataNode
[edureka@localhost sbin]$
```

Fig: Hadoop Installation – Starting ResourceManager

Start NodeManager:

The NodeManager in each machine framework is the agent which is responsible for managing containers, monitoring their resource usage and reporting the same to the ResourceManager.

Command: ./yarn-daemon.sh start nodemanager



```
edureka@localhost:~/hadoop-2.7.3/sbin
File Edit View Search Terminal Help
[edureka@localhost sbin]$ ./yarn-daemon.sh start nodemanager
starting nodemanager, logging to /home/edureka/hadoop-2.7.3/logs/yarn-edureka-nodemanager-localhost.localdomain.out
[edureka@localhost sbin]$ jps
22592 Jps
22113 NameNode
22310 ResourceManager
22206 DataNode
22559 NodeManager
[edureka@localhost sbin]$
```

Fig: Hadoop Installation – Starting NodeManager

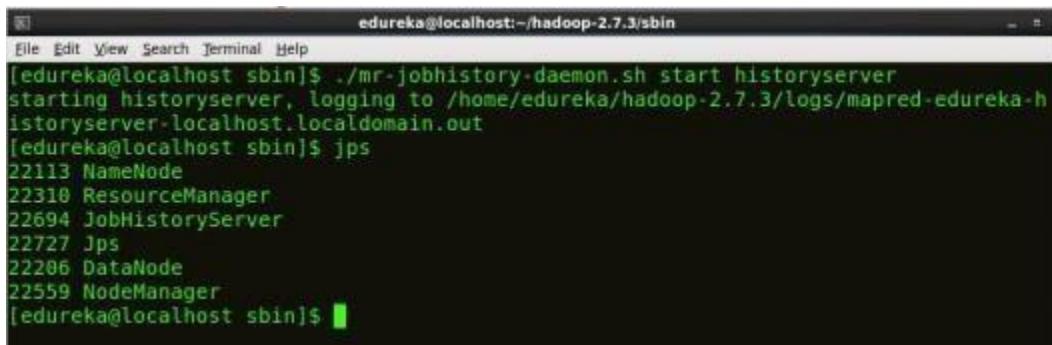
Start JobHistoryServer:

JobHistoryServer is responsible for servicing all job history related requests from client.

Command: ./mr-jobhistory-daemon.sh start historyserver

Step 14: To check that all the Hadoop services are up and running, run the below command.

Command: jps



```
edureka@localhost:~/hadoop-2.7.3/sbin
File Edit View Search Terminal Help
[edureka@localhost sbin]$ ./mr-jobhistory-daemon.sh start historyserver
starting historyserver, logging to /home/edureka/hadoop-2.7.3/logs/mapred-edureka-h
istoryserver-localhost.localdomain.out
[edureka@localhost sbin]$ jps
22113 NameNode
22318 ResourceManager
22694 JobHistoryServer
22727 Jps
22286 DataNode
22559 NodeManager
[edureka@localhost sbin]$
```

Fig: Hadoop Installation – Checking Daemons

Step 15: Now open the Mozilla browser and go to **localhost:50070/dfshealth.html** to check the NameNode interface.

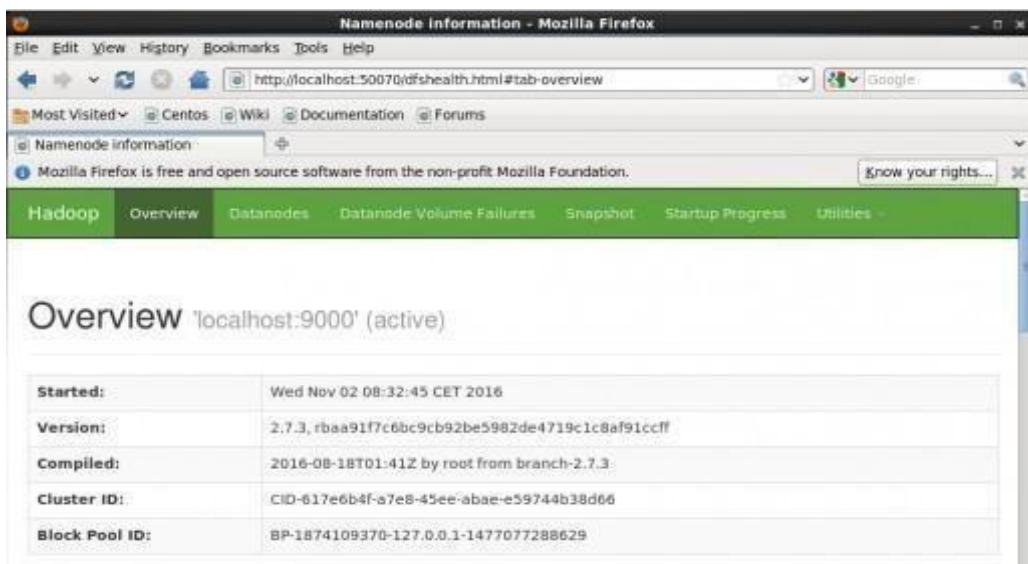


Fig: Hadoop Installation – Starting WebUI

- Congratulations, you have successfully installed a single node Hadoop cluster

Result:

EXP NO 7B)
DATE:

**WORD COUNT PROGRAM TO DEMONSTRATE THE USE OF MAP AND
REDUCE TASKS**

Aim:

Procedure:

step1

```
hduser@nspublin:/usr/local/hadoop/sbin$ mkdir /home/hduser/wc
```

step2:Compiling the java file - WordCount.java

```
hduser@nspublin:/usr/local/hadoop/sbin$ sudo /usr/lib/jvm/java-8-oracle/bin/javac -classpath  
/home/hduser/hadoop-core-1.2.1.jar -d /home/hduser/wc /home/hduser/WordCount.java
```

step3:Creating jar file for wordCount.java:

```
hduser@nspublin:/usr/local/hadoop/sbin$ jar -cvf /home/hduser/wc.jar -C /home/hduser/wc/ .
```

```
added manifest
```

```
adding: WordCount$IntSumReducer.class(in = 1739) (out= 739)(deflated 57%)
```

```
adding: WordCount$TokenizerMapper.class(in = 1736) (out= 753)(deflated 56%)
```

```
adding: WordCount.class(in = 1491) (out= 814)(deflated 45%)
```

step4.Executing jar file for WordCount.java

```
hduser@ksrietcsevb:/usr/local/hadoop/sbin$ hadoop jar /home/hduser/wc.jar WordCount /user/input  
/user/output
```

```
Picked up JAVA_TOOL_OPTIONS: -javaagent:/usr/share/java/jayatanaag.jar
```

```
16/09/12 10:52:53 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform...  
using builtin-java classes where applicable
```

Step5: to check the file in Output file

```
hduser@ksrietcsevb:/usr/local/hadoop/sbin$ hadoop fs -ls /user/output
```

```
Picked up JAVA_TOOL_OPTIONS: -javaagent:/usr/share/java/jayatanaag.jar
```

```
16/09/12 10:56:22 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform...  
using builtin-java classes where applicable
```

```
-rw-r--r-- 1 hduserssupergroup      0 2016-09-12 10:56 /user/output/_SUCCESS  
-rw-r--r-- 1 hduserssupergroup    182 2016-09-12 10:56 /user/output/part-r-00000
```

Program:

/home/hduser/WordCount.java:

```
import java.io.IOException; import  
java.util.StringTokenizer;  
import org.apache.hadoop.conf.Configuration; import  
org.apache.hadoop.fs.Path;  
import org.apache.hadoop.io.IntWritable; import  
org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Job; import  
org.apache.hadoop.mapreduce.Mapper; import  
org.apache.hadoop.mapreduce.Reducer;  
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat; import  
org.apache.hadoop.mapreduce.lib.output.FileOutputFormat; public class  
WordCount {  
    public static class TokenizerMapper  
        extends Mapper<Object, Text, Text, IntWritable>{ private final static  
        IntWritable one = new IntWritable(1); private Text word = new  
        Text();  
        public void map(Object key, Text value, Context context ) throws  
        IOException, InterruptedException {  
            StringTokenizeritr = new StringTokenizer(value.toString());  
            while (itr.hasMoreTokens()) {  
                word.set(itr.nextToken()); context.write(word,  
                one); } }  
    }  
    public static class IntSumReducer  
        extends Reducer<Text,IntWritable,Text,IntWritable> { private  
        IntWritable result = new IntWritable();  
        public void reduce(Text key, Iterable<IntWritable>values,Context context) throws IOException,
```

```

InterruptedException { int sum= 0;
for (IntWritableval : values) { sum += val.get();
} result.set(sum); context.write(key, result); } }
public static void main(String[] args) throws Exception {
Configuration conf = new Configuration();
Job job = Job.getInstance(conf, "word count");
job.setJarByClass(WordCount.class);
job.setMapperClass(TokenizerMapper.class);
job.setCombinerClass(IntSumReducer.class);
job.setReducerClass(IntSumReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
System.exit(job.waitForCompletion(true) ? 0 : 1); }
}

```

Input:

hai i am in gcc lab

Output

```

hduser@ksrietcsevb:/usr/local/hadoop/sbin$ hadoop fs -cat /user/output/*
Picked up JAVA_TOOL_OPTIONS: -javaagent:/usr/share/java/jayatanaag.jar
16/09/12 10:56:32 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform...
using builtin-java classes where applicable

```

am	1
gcc	1
hai	1
i	1
in	1
lab	1

hduser@ksrietcsevb:/usr/local/hadoop/sbin\$

Result:

EXP NO: 8	CREATING AND EXECUTING YOUR FIRST CONTAINER
DATE:	USING DOCKER

Aim:

Description: Creating and executing your first container using Docker is a fundamental step in understanding containerization. Here's a step-by-step guide to help you through the process.

Steps:**Step 1: Install Docker**

First, you need to install Docker on your system. Docker provides installers for various operating systems. You can find the installation instructions on the Docker website.

Step 2: Verify Installation

After installing Docker, verify that it's properly installed by running the following command in your terminal or command prompt:

```
Copy code
css
docker --version
```

This command should display the Docker version if it has been installed correctly.

Step 3: Pull a Docker Image

Docker containers are created based on Docker images. To create your first container, you need to pull an image from Docker Hub. Docker Hub is a repository of Docker images maintained by Docker. As an example, let's pull the official hello-world image, which is a simple image that prints a message:

```
Copy code
docker pull hello-world
```

This command downloads the hello-world image from Docker Hub to your local machine.

Step 4: Run a Docker Container

Once you have pulled the image, you can run a container based on that image using the following command:

```
arduino
```

 Copy code

```
docker run hello-world
```

When you run this command, Docker executes the hello-world container. The container will print a "Hello from Docker!" message, indicating that your installation appears to be working correctly.

Step 5: Explore Docker Containers

To see the list of running containers, you can use:

```
 Copy code
```

```
docker ps
```

This command shows all running containers. Since the hello-world container completes its execution quickly, you might not see it in the list if you run docker ps after it has finished running. To see all containers, including the ones that have exited, use:

```
css
```

 Copy code

```
docker ps -a
```

Step 6: Cleanup (Optional)

After you've run your first container, you may want to clean up unused containers and images from your system to free up space. Use the following commands to remove containers and images:

- To remove a specific container:

```
bash
```

 Copy code

```
docker rm CONTAINER_ID
```

To remove a specific image:

```
 Copy code
```

```
docker rmi IMAGE_ID
```

To remove all stopped containers:

```
 Copy code
```

```
docker container prune
```

To remove all unused images

```
arduino
```

```
 Copy code
```

```
docker image prune
```

Hints:

- **Docker Images:** Explore more Docker images on Docker Hub (hub.docker.com). There are images for various applications and services that you can use to create containers.
- **Docker file:** For more complex setups, learn about Docker files, which are text files that contain instructions to build a Docker image automatically.

By following these steps, you have successfully created and executed your first Docker container. Docker provides a powerful way to package and deploy applications in a consistent and reproducible environment. Experiment with different images and explore Docker's capabilities to further enhance your understanding of containerization.

Result:

EXP NO: 9 DATE:	RUN A CONTAINER FROM DOCKER HUB
----------------------------------	--

Aim:

Description: Running a container from Docker Hub is a straightforward process using the Docker command-line interface (CLI). Docker Hub is a repository of Docker images maintained by Docker, Inc., and the community. Here's a step-by-step guide to running a container from Docker Hub

Steps:**1. Install Docker:**

Ensure Docker is installed on your system. You can download Docker Desktop for Windows or Mac from Docker's official website or install Docker Engine on Linux by following the instructions for your specific distribution.

2. Search for an Image:

- If you haven't decided on an image yet, you can explore Docker Hub at hub.docker.com. Docker Hub provides a search bar where you can find various Docker images for different applications and services.

3. Pull the Docker Image:

- Open your terminal or command prompt.
- Use the docker pull command to download the desired Docker image from Docker Hub. Replace image_name with the name of the image you want to pull. For example,to pull the official Nginx web server image:

```
Copy code
docker pull nginx
```

If you want a specific version of an image, you can specify the version tag. For example

```
Copy code
docker pull nginx:1.21.1
```

4. Run the Docker Container:

- Once the image is pulled, you can run a container based on that image using the docker run command. Replace image_name with the image you pulled earlier. For example, to run a container from the Nginx image:

```
arduino                               Copy code
docker run -d -p 80:80 nginx
```

- d flag runs the container in detached mode (in the background).
- p 80:80 publishes port 80 of the container to port 80 on the host (your local machine). This allows you to access the Nginx web server running inside the container from your web browser.
- nginx is the name of the image you pulled.

5. Verify the Container is Running:

To see the list of running containers, use the docker ps command.

```
Copy code
docker ps
```

- This command lists all containers that are currently running. You should see the Nginx container (or any other container you started) listed here.

6. Access the Containerized Application:

- Open a web browser and go to http://localhost or http://127.0.0.1 (depending on your operating system) to see the default Nginx welcome page if you ran an Nginx container. If you mapped a different port, use http://localhost:your_port.

7. Interact with the Container (Optional):

- You can interact with a running container using various Docker CLI commands, such as docker exec to execute commands inside the container, docker logs to view container logs, etc.

8. Stop and Remove the Container (Optional):

- To stop a running container, use docker stop followed by the container ID or name:

```
arduino
```

 Copy code

```
docker stop container_id_or_name
```

To remove a stopped container, use docker rm followed by the container ID or name

```
bash
```

 Copy code

```
docker rm container_id_or_name
```



Hints:

- Docker Hub provides a vast repository of images for different applications and services, including databases, web servers, development environments, and more.
- Make sure to review the documentation or README file associated with the Docker image you're pulling for specific instructions or configurations.
- Always consider security implications and configure container security appropriately, especially when running containers from third-party sources.

By following these steps, you can easily run a container from Docker Hub and start utilizing various applications and services in a containerized environment on your local machine.

Result:

EXP NO 10 DATE:	CREATE A VIRTUALIZATION IN VMWARE OR ANY EQUIVALENT OPEN SOURCE TOOL TO ALLOCATE MEMORY AND STORAGE SPACE AS PER REQUIREMENT AND INSTALL GUEST OS ON THAT VMWARE
----------------------------	---

Aim:

Description:

Creating a virtual machine (VM) using VMware or an equivalent open-source tool involves several steps, including setting up the VM configuration, allocating memory and storage, and installing a guest operating system (OS). Here's a detailed guide using VMware Workstation as an example.

Steps:

Step 1: Install VMware Workstation

1. **Download VMware Workstation:**
 - Go to the VMware website and download the latest version of VMware Workstation.
2. **Install VMware Workstation:**
 - Run the installer and follow the on-screen instructions to complete the installation.

Step 2: Create a New Virtual Machine

1. **Open VMware Workstation:**
 - Launch VMware Workstation from your desktop or start menu.
2. **Start the New VM Wizard:**
 - Click on "Create a New Virtual Machine" or go to File > New Virtual Machine.
3. **Choose Configuration Type:**
 - Select "Typical (recommended)" for the configuration type and click "Next".

Step 3: Select the Installation Media

1. **Choose Installation Media:**
 - You have three options:
 - Install from a disc.
 - Install from an ISO image file.
 - Install later.
 - Select the appropriate option based on your installation media and click "Next".

Step 4: Select the Guest Operating System

1. Choose OS Type:

- Select the guest operating system you plan to install (e.g., Windows, Linux) and the version. Click “Next”.

Step 5: Name the Virtual Machine

1. Name Your VM:

- Enter a name for your virtual machine and specify the location where you want to store the VM files. Click “Next”.

Step 6: Specify Disk Capacity

1. Set Disk Size:

- Specify the maximum disk size for the VM. You can also choose whether to store the virtual disk as a single file or split it into multiple files. Click “Next”.

Step 7: Customize Hardware (Optional)

1. Customize Hardware Settings:

- You can customize the VM hardware settings by clicking on “Customize Hardware”. Here, you can adjust the amount of memory, number of processors, network settings, and more.
- After customizing, click “Close” to return to the main wizard.

Step 8: Complete the VM Creation

1. Finish:

- Review your settings and click “Finish” to create the virtual machine.

Step 9: Install the Guest Operating System

1. Start the VM:

- Select your newly created VM and click on “Power on this virtual machine”.

2. Follow OS Installation Process:

- Follow the on-screen instructions to install the guest operating system just as you would on a physical machine.

Step 10: Install VMware Tools

1. Install VMware Tools:

- After the guest OS installation is complete, it's recommended to install VMware Tools. This improves the VM's performance and enhances its usability.

- To install VMware Tools, go to VM > Install VMware Tools and follow the instructions within the guest OS.

Hints:

- **Snapshots:** Take snapshots of your VM at different points in time to easily revert to a previous state if needed. This can be done from the VM > Snapshot menu.
- **Backups:** Regularly back up your VM files to prevent data loss.
- **Updates:** Keep both VMware Workstation and VMware Tools up to date for the best performance and security.

By following these steps, you can create and manage a virtual machine using VMware Workstation effectively. If you prefer using an open-source alternative, you can follow similar steps with tools like Virtual Box or KVM.

Result:

EXP NO: 11
DATE:

TO HOST A WEB PAGE IN EC2 USING AWS.

Aim:

Description: Hosting a web page on an Amazon EC2 instance involves several steps. Here's a detailed guide to help you set up and deploy a simple web page using AWS EC2:

Steps:

Step 1: Set Up an AWS Account

1. Create an AWS Account:

- Go to the [AWS Management Console](#) and sign up for a new account if you don't already have one.
- Follow the on-screen instructions to complete the registration process.

Step 2: Launch an EC2 Instance

1. Open EC2 Dashboard:

- Sign in to the AWS Management Console.
- Navigate to the EC2 Dashboard by searching for "EC2" in the services search bar and selecting it.

2. Launch an Instance:

- Click on the "Launch Instance" button.
- Select an Amazon Machine Image (AMI). For this guide, you can choose the Amazon Linux 2 AMI (Free Tier eligible).

3. Choose an Instance Type:

- Select an instance type. The t2.micro instance type is free-tier eligible and sufficient for hosting a simple web page. Click "Next: Configure Instance Details".

4. Configure Instance Details:

- You can accept the default settings here. Click "Next: Add Storage".

5. Add Storage:

- The default 8 GB storage is sufficient. Click "Next: Add Tags".

6. Add Tags:

- Add a tag to name your instance (optional). For example, you can use Key: Name, Value: MyWebServer. Click "Next: Configure Security Group".

7. Configure Security Group:

- Create a new security group. Add rules to allow HTTP (port 80) and SSH (port 22) traffic:
 - Type: HTTP, Protocol: TCP, Port Range: 80, Source: Anywhere (0.0.0.0/0)
 - Type: SSH, Protocol: TCP, Port Range: 22, Source: Anywhere (0.0.0.0/0)
- Click "Review and Launch".

8. Review and Launch:

- Review your instance settings. Click "Launch".

- When prompted, create a new key pair (or use an existing one). Download the key pair file (e.g., my-key-pair.pem) and keep it secure. Click “Launch Instances”.

Step 3: Connect to Your EC2 Instance

1. Access Your Instance:

- Go to the “Instances” section in the EC2 Dashboard.
- Select your instance and click on the “Connect” button.
- Follow the instructions to connect to your instance using SSH. Typically, you will use a terminal or an SSH client with the following command:

```
sh
ssh -i "my-key-pair.pem" ec2-user@your-ec2-public-dns
Copy code
```

Step 4: Set Up a Web Server

1. Update Your Instance:

- Once connected, update the package list

2. Install Apache Web Server:

- Install Apache (httpd)

```
sh
sudo yum install -y httpd
Copy code
```

Start Apache:

Start the Apache service

```
sh
sudo systemctl start httpd
Copy code
```

Enable Apache to start on boot:

```
sh
sudo systemctl enable httpd
Copy code
```

Step 5: Deploy Your Web Page

1. Create a Simple Web Page:

- Navigate to the Apache default root directory:

```
sh
```

[Copy code](#)

```
cd /var/www/html
```

- Create an `index.html` file:

```
sh
```

[Copy code](#)

```
sudo nano index.html
```

- Add the following HTML content (or your own):

```
html
```

[Copy code](#)

```
<!DOCTYPE html>  
<html lang="en">
```



- Add the following HTML content (or your own):

```
html
```

[Copy code](#)

```
<!DOCTYPE html>  
<html lang="en">  
  <head>  
    <meta charset="UTF-8">  
    <meta name="viewport" content="width=device-width, initial-scale=1.0">  
    <title>Welcome to My Web Page</title>  
  </head>  
  <body>  
    <h1>Hello, World!</h1>  
    <p>Welcome to my web page hosted on AWS EC2.</p>  
  </body>  
</html>
```



- Save and exit the file (Ctrl+X, Y, Enter nano).

Step 6: Access Your Web Page

1. Open Your Web Page:

- Open a web browser and enter your EC2 instance's public DNS or public IP address (found in the EC2 Dashboard).
- You should see your web page with the content you added to index.html.

Step 7: Manage and Secure Your Instance

1. Monitor Your Instance:

- Use the EC2 Dashboard to monitor the status and performance of your instance.

2. Secure Your Instance:

- Regularly update your instance.
- Use a strong password for SSH access or consider using key pairs exclusively.
- Configure a firewall and other security best practices as needed.

By following these steps, you can successfully host a web page on an EC2 instance using AWS.

Result:

EXP NO: 12
DATE:

TO CREATE BILLING ALARM USING AWS ACCOUNT.

Aim:

Description: Creating a billing alarm in your AWS account helps you monitor your costs and ensure that you don't exceed your budget. Here's a step-by-step guide to set up a billing alarm using AWS Cloud Watch.

Step 1: Sign In to the AWS Management Console

1. Go to the AWS Management Console:

- Open your web browser and navigate to the [AWS Management Console](#).
- Sign in with your AWS account credentials.

Step 2: Open the Cloud Watch Console

1. Navigate to Cloud Watch:

- In the AWS Management Console, type "CloudWatch" in the search bar and select it from the list of services.

Step 3: Access the Billing Section

1. Go to the Billing Dashboard:

- In the Cloud Watch console, on the left navigation pane, expand "Alarms" and click on "Billing".
- If you don't see the billing option, ensure your AWS account is using a region that supports billing alarms (preferably the US East (N. Virginia) region).

Step 4: Create a Billing Alarm

1. Create Alarm:

- Click the "Create Alarm" button.

Step 5: Specify Alarm Conditions

1. Specify Metric and Conditions:

- In the "Create Alarm" wizard, under "Select Metric", click on the "Select metric" button.
- In the "Browse" tab, choose "Billing" and then select "Total Estimated Charge".
- Click on "Select metric" to proceed.

2. Set Conditions:

- In the "Specify metric and conditions" step, set the "Threshold type" to "Static".
- Define the threshold value for your alarm (e.g., 100 USD). This is the amount at which you want to be notified.
- Choose the "Whenever the estimated charges is..." option and set it to "Greater/Equal" to your threshold value.

Step 6: Configure Actions

1. Set Notification:

- In the “Configure actions” step, choose to “Select an existing SNS topic” or “Create a new topic”.
- If creating a new topic, provide a topic name and enter the email addresses that should receive the notifications. Confirm the subscription by checking your email and clicking the confirmation link sent by AWS.
- Select the SNS topic you created or chose for notification.

Step 7: Add a Name and Description

1. Name Your Alarm:

- Provide a name and a description for your alarm in the “Add name and description” step (e.g., “Monthly Billing Alarm”).
- Click “Next” to review your alarm settings.

Step 8: Review and Create the Alarm

1. Review Your Settings:

- Review all the settings you have configured for the billing alarm.
- Click “Create Alarm” to finish the setup.

Step 9: Confirm Email Subscription

1. Confirm Subscription:

- If you created a new SNS topic and added email addresses, make sure the recipients confirm their subscriptions by clicking the link in the email sent from AWS Notifications.

Step 10: Monitor Your Alarm

1. Monitor the Alarm:

- The alarm will now be listed under the “Alarms” section in CloudWatch.
- You will receive an email notification when your AWS charges exceed the threshold you set.

Additional Tips:

- **Budget Reports:** For more detailed cost management, consider using AWS Budgets to track your costs and usage against your budget.
- **Threshold Adjustment:** You can adjust the threshold or settings of your billing alarm anytime from the Cloud Watch console.
- **Multiple Alarms:** Create multiple billing alarms if you need to monitor different services or different thresholds.

By following these steps, you can create and manage billing alarms to keep track of your AWS spending effectively.

Result

CONTENT BEYOND THE SYLLABUS

EX.NO: 13	<u>WORDMAPPER PROGRAM TO DEMONSTRATE THE</u>
DATE:	<u>USE OF WINDOWS 7 OR 8.</u>

AIM :

To Develop Word Mapper program for demonstrating the use of Map and Reduce tasks

PROCEDURE :

1. Prepare the example input CSV file with English words from a dictionary and all translations in other languages added to it, separated by a ‘|’ symbol.
2. Create the Java Map Reduce project
3. Implement the Map Reduce classes
 - i) The first class is the “Mapper” : Word Mapper
 - ii) The next class is the “Reducer” which reduces the map to the wanted output.
: All Translations Reducer
- This Reduce steps collects all values for a given key and put them after each other separated with a ‘|’ symbol.
4. The applications using Java by using Grid API is successfully complete

SAMPLE OUTPUT:

Result:

Thus the word mapper program was successfully implemented

CONTENT BEYOND THE SYLLABUS

EX.NO: 14	<u>CLOUD STORAGE IN DROP BOX</u>
DATE:	

The cloud storage provider and simple ways you can start interacting with it on your desktop and easy-to-use low cost cloud storage platform.



Dropbox is a cloud storage service, which means you can copy your files to the cloud and access them later, even if you're using a different device.

DropboxSyncing

That doesn't free space on your hard drive, though. A copy of the file you save to Dropbox will remain on your hard drive.

If you change the file locally, it will be updated in the cloud.

It's unfortunate that Dropbox doesn't free up hard drive space, at least on its free plan. That comes with an upside, though. Dropbox syncs your data across all devices. As long as you can access your Dropbox account, you can download any file stored in your Dropbox folder on your local machine.

This system keeps files secure against technical problems you may have with your computer. For instance, if your hard drive fails, your photos and other important documents won't be lost forever if they're stored in the cloud. Although you may think your computer is secure, there's always a small chance it will be damaged, die of its own accord or even get stolen.

Dropbox makes it easy to transfer those important files, photos and folders to a new computer, which saves you a lot of time and frustration, as well.

It also makes it easy if you want to share a file with colleagues and family members. In addition to sharing, you can also use Dropbox to view files on multiple devices. For example, you can move pictures you took on your phone to Dropbox and look at the images on your computer later.

You can use Dropbox as pure cloud storage, too. As long as you use the [web client](#), you can add files to your Dropbox account and store them only in the cloud.

Dropbox Pricing & Account Tiers

Basic only comes with 2GB of free storage, which is enough to start using Dropbox, but not much else. We'll show you how to sign up for the free Dropbox account in the next section.

Dropbox has multiple paid tiers for personal accounts, as well.

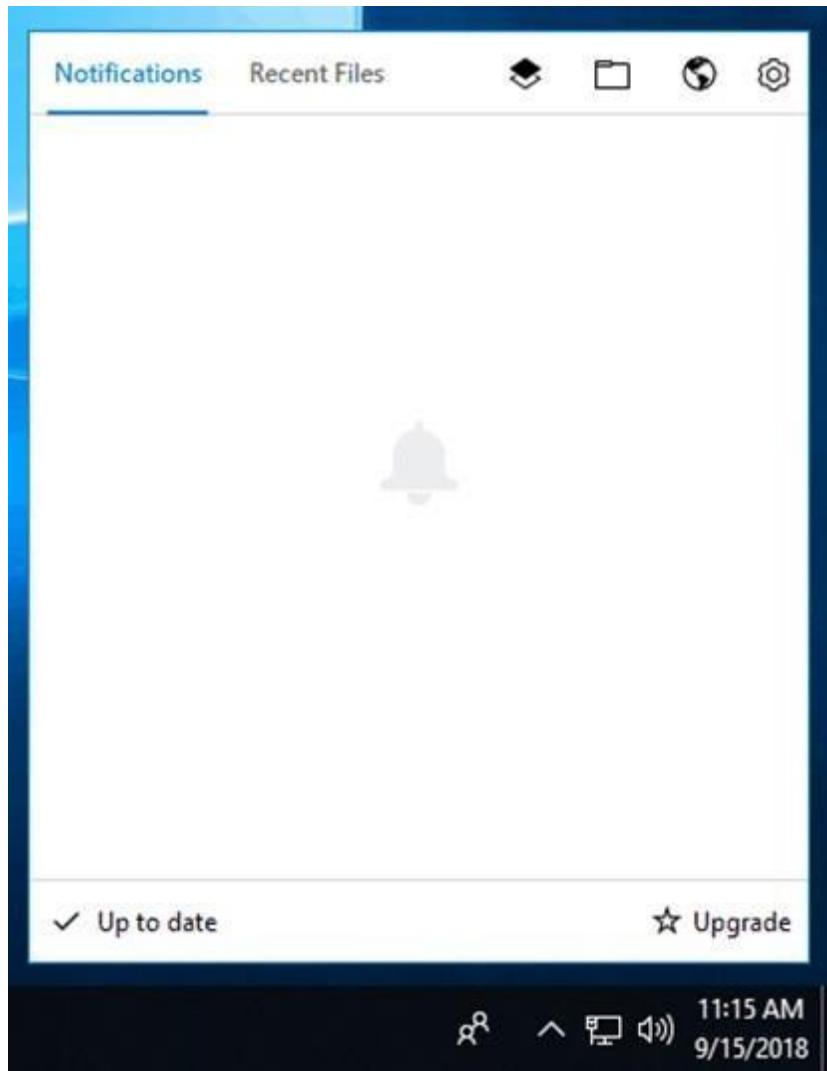
Setting Up Dropbox

One of the first things you will notice after you've installed Dropbox is that it added an icon to your system tray. The tray icon is the easiest and fastest way to access Dropbox from your computer.

Left-click once or right-click to open Dropbox.

There are two tabs in the upper-left corner and four buttons in the upper right.

The first tab is "notifications," which will show you the notifications you want, as well as occasional messages from Dropbox.



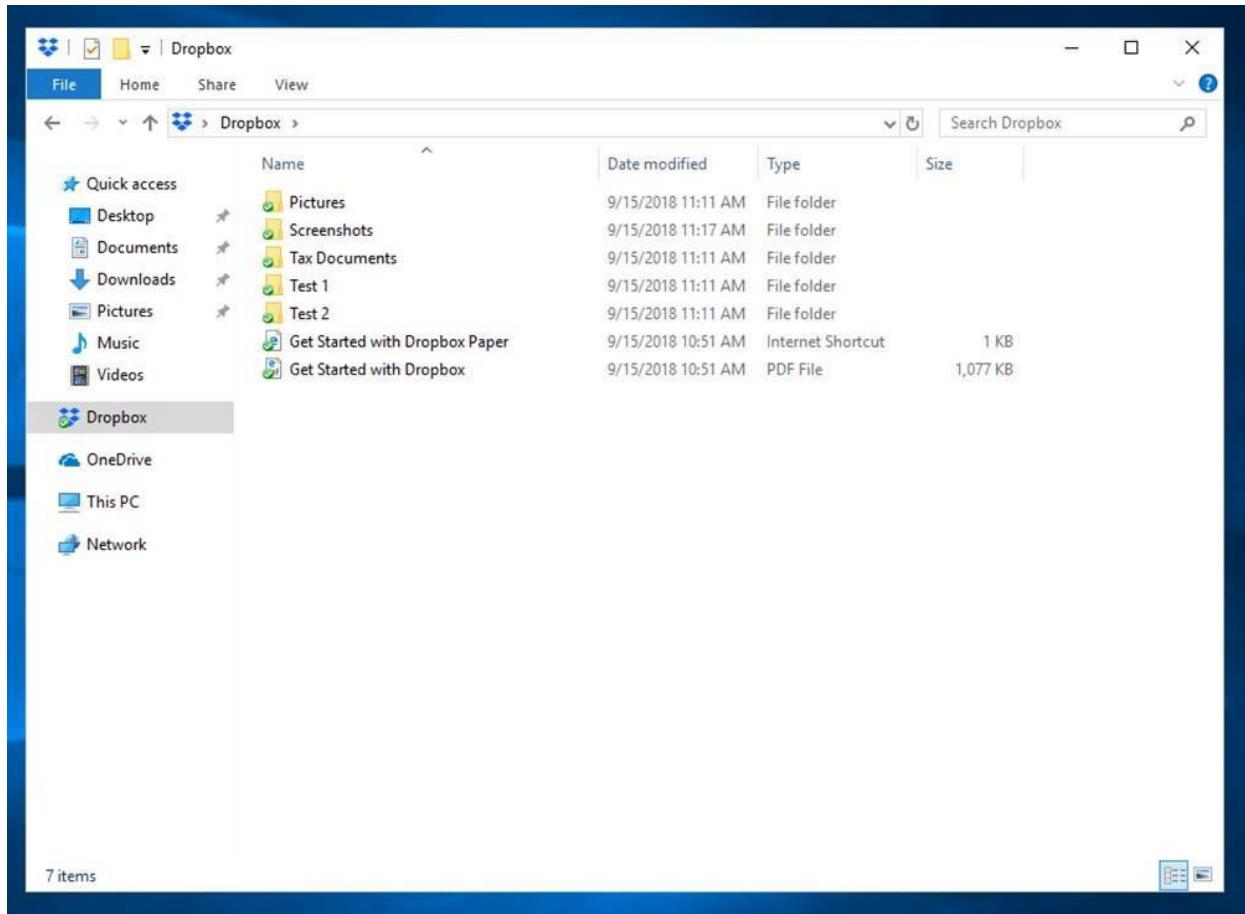
How to Use Dropbox on Desktop

Dropbox on your desktop is easy. Downloading the client adds a folder under your username titled “Dropbox,” and the files you add to that folder are automatically copied to your Dropbox account.

You can find the folder by opening File Explorer and navigating to the section called “This PC” on Windows. Then, you’ll double-click on the drive you installed Dropbox on. That will usually be your C: drive.

Next, double-click on the folder called “Users,” then the folder with your username. In that folder, you’ll find Dropbox. Double-clicking on its icon or name opens the Dropbox folder. This folder is the same one you access when you click the second button in the icon menu mentioned above.

In Windows, you will likely also find the Dropbox folder under “quick access.”



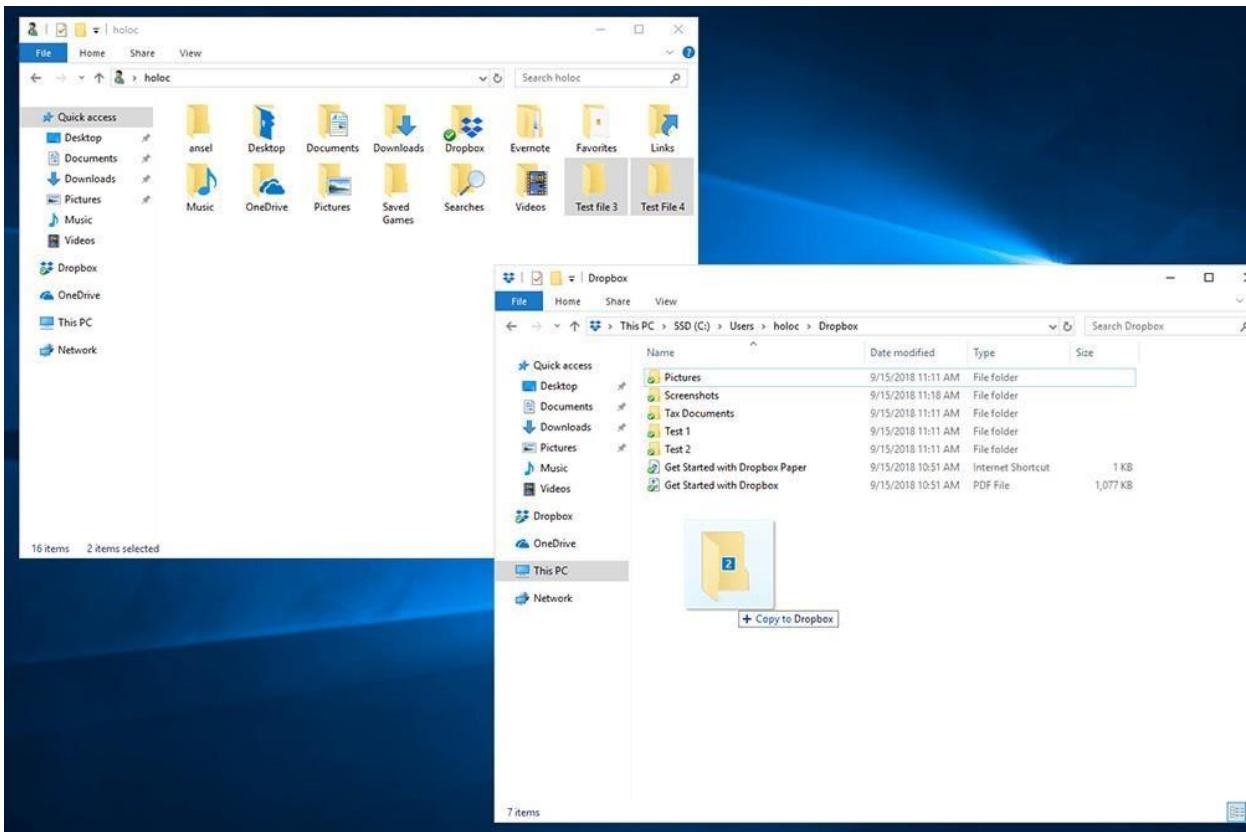
Transferring Files With Dropbox

If you want to move old files to Dropbox, navigate to them in a second window. Once you've found them, simply highlight the file you want to move, then click and drag it over to your Dropbox folder.

You can also copy and paste files or folders to Dropbox if you find that easier.

It may take a second or two to transfer files to the Dropbox folder. Once they're done, Dropbox will show a blue icon in the lower left with arrows on it. The icon means it is copying the file and uploading it to the cloud. Once it's done, the icon will turn green and display a checkmark.

If you wish to transfer multiple files at once, hold the "ctrl" button and click on each file you want to move. You'll know you've done it right if the files you click on remain highlighted. Then, click and drag one of the files to move all those selected into the Dropbox folder.



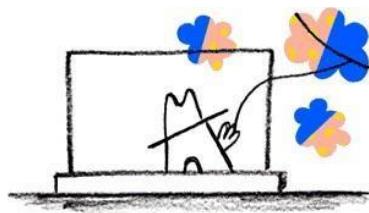
How to Use Dropbox on Mobile and Web

Using Dropbox from the web client is a similar experience to using it on the desktop. When you log in to your Dropbox account, you're presented with the files recently changed on your desktop.

It is the same list you'd find by clicking on the “recent files” tab in the icon menu.

Download Dropbox to get started

[Download Dropbox](#)



SAMPLE OUTPUT:

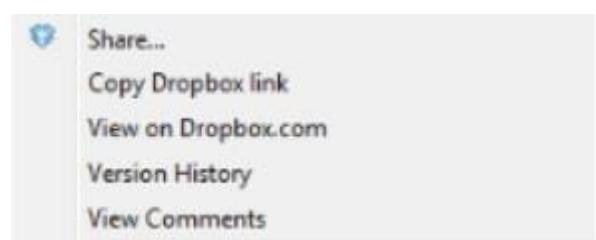
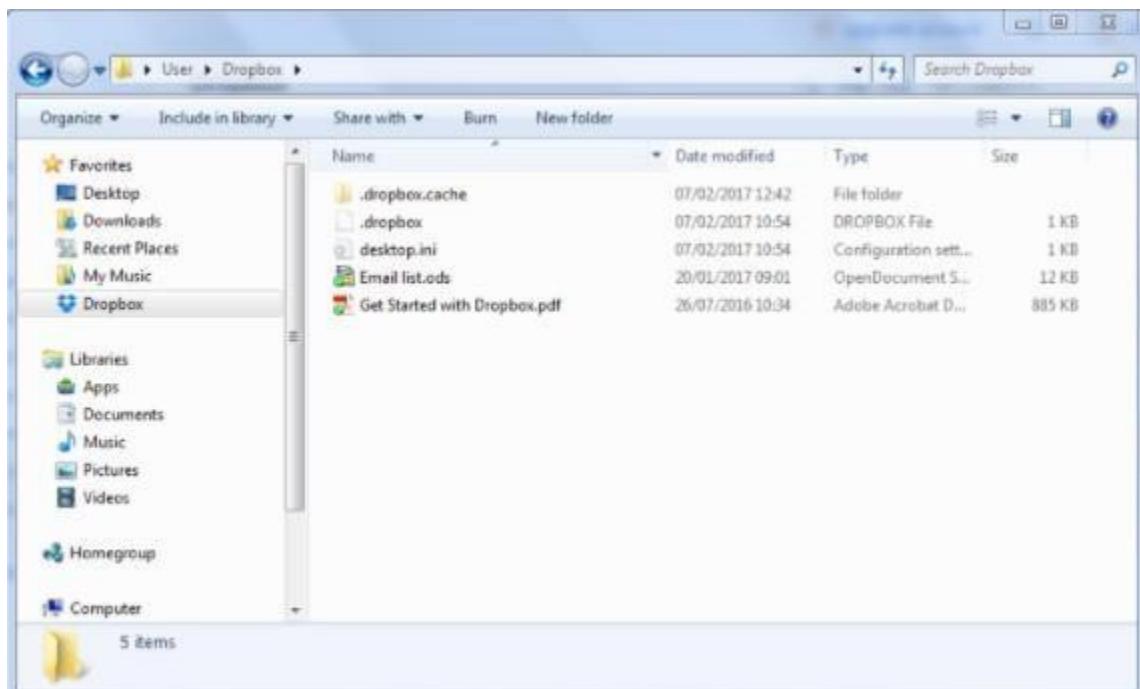
Dropbox

Search

Name	Modified	Members	More
Pictures	--	Only you	...
Screenshots	--	Only you	Share ...
Tax Documents	--	Only you	...
Test 1	--	Only you	...
Test 2	--	Only you	...

Add to Starred





Result:

Thus the usage of cloud storage in drop box is learned successfully